

# all\_models

December 4, 2022

## 0.1 Imports

```
[ ]: import pandas as pd
import numpy as np
import tensorflow as tf
import os

import seaborn as sns
import matplotlib.pyplot as plt
import sys

import skimage
from skimage.color import rgb2hsv
from skimage.transform import rescale, resize
from tqdm import tqdm

import sys
import os

from sklearn.metrics import classification_report

import bz2, pickle, _pickle as cPickle

import random

# random.seed(1234)

# module_path = os.path.abspath(os.path.join '..', '..'))
# if module_path not in sys.path:
#     sys.path.append(module_path+"/Modules/Testing")
# import testing_module

SAVE_DIR = "../Pickled Datasets/"

COPIES = 2
N_DIGITS = 3
HEIGHT = 25
```

```

WIDTH = 25

PLOT_SAVE = "../Plots/loss-curves/"

def compressed_pickle(name: str, data):
    with bz2.BZ2File(os.path.join(SAVE_DIR, "{}.pbz2".format(name)), 'w') as f:
        cPickle.dump(data, f)

def decompress_pickle(file):
    data = bz2.BZ2File(file, 'rb')
    data = cPickle.load(data)
    return data

def plot_history(history):
    acc=history.history['accuracy']
    val_acc=history.history['val_accuracy']
    loss=history.history['loss']
    val_loss=history.history['val_loss']
    epochs=range(len(acc))

    fig, ax = plt.subplots(1, 2, figsize = (12, 6))
    ax[0].plot(epochs, acc, 'r', label = "Training Accuracy")
    ax[0].plot(epochs, val_acc, 'b', label = "Validation Accuracy")
    ax[0].legend()
    ax[0].set_title('Training and Validation Accuracy')
    ax[0].set_xlabel("Epochs")

    ax[1].plot(epochs, loss, 'r', label = "Training Loss")
    ax[1].plot(epochs, val_loss, 'b', label = "Validation Loss")
    ax[1].set_title('Training and Validation Losses')
    ax[1].set_xlabel("Epochs")
    plt.show()

    return (fig, ax)

```

SKImage rescales the image for us! Which means that we don't need to rescale by 255.0 anymore, saving us needlessly spent time and effort. There is another Augmentor library which can be used for data augmentation. We can simply sample the augmented images henceforth!

## 0.2 Preliminary setup

```

[ ]: # sys.path.append(os.path.dirname(os.path.join((os.path.pardir), "Modules")))

# origin_dir = os.path.join(os.path.pardir, 'Data')
# new_dir_path = os.path.join(os.path.pardir, 'Data', 'cell_images')

```

```

# #for local systems
# train_csv = os.path.join(origin_dir, 'train.csv')
# test_csv = os.path.join(origin_dir, 'test.csv')
# val_csv = os.path.join(origin_dir, 'val.csv')

# from Modules.labelling import Labelling

# # download = Data_Download(origin_dir)
# # data_dir = download.resize_image(new_dir_path, 44, 44)

# lab = Labelling()
# lab.label('../Data/cell_images/', exclude_mislabeled= True)      # function
#                               ↪to label the dataset
# train_csv, val_csv, test_csv = lab.train_test_val_split('../Data/', '../Data/
#                               ↪cell_images/labels.csv', random_state = 1234)

# train_data = pd.read_csv(train_csv)
# val_data    = pd.read_csv(val_csv)
# test_data   = pd.read_csv(test_csv)

```

## 0.2.1 Reading images

```

[ ]: # def read_image(path):
#     '''Function to read images given a path and return an array'''
#     return skimage.io.imread(path)

# i = 14

# print(train_data['Image_Path'][i])
# image = rgb2hsv(skimage.io.imread(train_data['Image_Path'][i]))
# print(np.max(image))
# result = ((image > 0.5)*image)[..., 1]
# plt.imshow(result, 'gray')

# tqdm.pandas()
# train_data['image_arr'] = train_data['Image_Path'].progress_apply(lambda x:
#                               ↪read_image(x))
# val_data['image_arr']     = val_data['Image_Path'].progress_apply(lambda x:
#                               ↪read_image(x))
# test_data['image_arr']    = test_data['Image_Path'].progress_apply(lambda x:
#                               ↪read_image(x))

# x_train, y_train = train_data['image_arr'].to_numpy(),
#                               ↪train_data['Parasitized'].to_numpy()

```

```
# x_val , y_val = val_data['image_arr'].to_numpy() ,  
    ↪ val_data['Parasitized'].to_numpy()  
# x_test , y_test = test_data['image_arr'].to_numpy() ,  
    ↪ test_data['Parasitized'].to_numpy()
```

### 0.3 Data Augmentation

```
[ ]: # import albumentations as A  
# import cv2  
  
# augment = A.augmentations.geometric.transforms.Affine(  
#     translate_percent = 0.1,  
#     rotate = 60,  
#     shear = 30  
# )  
  
# augment = A.ShiftScaleRotate(scale_limit = (-0.2, 0), rotate_limit= 90,  
    ↪ border_mode=cv2.BORDER_CONSTANT, always_apply= True)  
  
# transform = A.Compose(  
#     [  
#         A.Resize(HEIGHT, WIDTH, always_apply= True),  
#         A.Rotate(90, border_mode=cv2.BORDER_CONSTANT),  
#         A.VerticalFlip(p = 0.5),  
#         A.HorizontalFlip(p = 0.5),  
#         A.augmentations.geometric.Affine(shear = 7.5, mode=cv2.  
    ↪ BORDER_CONSTANT),  
#         A.GaussNoise(var_limit = (0.01, 0.05))  
#     ]  
# )  
  
# aug_dataset = []  
# aug_labels = []  
  
# for i, lab in tqdm(zip(x_train, y_train)):  
#     for _ in range(COPIES):  
#         aug_dataset.append(transform(image = i)['image'])  
#         aug_labels.append(lab)  
  
# x_train_aug = np.array(aug_dataset)  
# y_train_aug = np.array(aug_labels)  
  
# np.unique(y_train_aug, return_counts = True)
```

### 0.3.1 Resizing

```
[ ]: # temp = []
# for img in tqdm(x_train):
#     temp.append(resize(img, (HEIGHT, WIDTH)))
# x_train = np.array(temp)

# temp = []
# for img in tqdm(x_val):
#     temp.append(resize(img, (HEIGHT, WIDTH)))
# x_val = np.array(temp)

# temp = []
# for img in tqdm(x_test):
#     temp.append(resize(img, (HEIGHT, WIDTH)))
# x_test = np.array(temp)
```

### 0.4 Saving Data

```
[ ]: # compressed_pickle("x_train_aug", x_train_aug)
# compressed_pickle("y_train_aug", y_train_aug)
# compressed_pickle("x_train", x_train)
# compressed_pickle("y_train", y_train)
# compressed_pickle("x_val", x_val)
# compressed_pickle("y_val", y_val)
# compressed_pickle("x_test", x_test)
# compressed_pickle("y_test", y_test)

# n_aug_train = x_train_aug.shape[0]
# n_train      = x_train.shape[0]
# n_val        = x_val.shape[0]
# n_test       = x_test.shape[0]
```

### 0.5 Loading Data

```
[ ]: x_train_aug = decompress_pickle(SAVE_DIR + 'x_train_aug.pbz2')
y_train_aug = decompress_pickle(SAVE_DIR + 'y_train_aug.pbz2')
x_train = decompress_pickle(SAVE_DIR + 'x_train.pbz2')
y_train = decompress_pickle(SAVE_DIR + 'y_train.pbz2')
x_val = decompress_pickle(SAVE_DIR + 'x_val.pbz2')
y_val = decompress_pickle(SAVE_DIR + 'y_val.pbz2')
x_test = decompress_pickle(SAVE_DIR + 'x_test.pbz2')
y_test = decompress_pickle(SAVE_DIR + 'y_test.pbz2')

print("augmented: ", x_train_aug.shape, y_train_aug.shape)
```


```
print("train: ", x_train.shape, y_train.shape)
print("val: ", x_val.shape, y_val.shape)
print("test: ", x_test.shape, y_test.shape)
```

```
augmented: (41202, 25, 25, 3) (41202,)
train: (20601, 25, 25, 3) (20601,)
val: (2943, 25, 25, 3) (2943,)
test: (2617, 25, 25, 3) (2617,)
```

```
[ ]: n_aug_train = x_train_aug.shape[0]
     n_train     = x_train.shape[0]
     n_val       = x_val.shape[0]
     n_test      = x_test.shape[0]
```

```
[ ]: plt.imshow(x_train[2]), x_train[2].shape
```

```
[ ]: (<matplotlib.image.AxesImage at 0x1e575f68280>, (25, 25, 1))
```



## 0.6 RGB Modeling

### 0.6.1 Unaugmented

#### Naive Bayes

```
[ ]: from sklearn import naive_bayes

nb_cls = naive_bayes.GaussianNB()
nb_cls.fit(x_train.reshape(n_train, -1), y_train)
preds_train = nb_cls.predict(x_train.reshape(n_train, -1))
preds_val = nb_cls.predict(x_val.reshape(n_val, -1))
preds_test = nb_cls.predict(x_test.reshape(n_test, -1))

print("Training Classification Report: \n", classification_report(y_train,
↪preds_train, digits = N_DIGITS))
print("\nValidation Classification Report: \n", classification_report(y_val,
↪preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test,
↪preds_test, digits = N_DIGITS))

# eval = testing_module.ModelEvaluation(y_val,
#     nb_cls.predict(x_val.reshape(n_val, -1)),
#     model_reference_name = 'Naive Bayes',
#     model_type = 'classification',
#     plot_classification_metric = ['roc_auc']) # if classification
# eval.evaluate(evaluate_save= True, plots_show = True)
```

#### Training Classification Report:

	precision	recall	f1-score	support
0.0	0.611	0.746	0.672	10260
1.0	0.677	0.530	0.594	10341
accuracy			0.637	20601
macro avg	0.644	0.638	0.633	20601
weighted avg	0.644	0.637	0.633	20601

#### Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.620	0.761	0.683	1466
1.0	0.693	0.538	0.606	1477
accuracy			0.649	2943
macro avg	0.657	0.649	0.644	2943
weighted avg	0.657	0.649	0.644	2943

#### Testing Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0.0	0.613	0.766	0.681	1303
1.0	0.692	0.521	0.595	1314
accuracy			0.643	2617
macro avg	0.653	0.644	0.638	2617
weighted avg	0.653	0.643	0.638	2617

## Logistic Regression

```
[ ]: from sklearn.linear_model import LogisticRegression

logreg_cls = tf.keras.Sequential([tf.keras.layers.Dense(2, activation =_
    ↳ 'sigmoid', input_dim = HEIGHT * WIDTH * 3)])
logreg_cls.compile(optimizer = tf.optimizers.Adam(0.001), loss =_
    ↳ 'sparse_categorical_crossentropy', metrics = ['accuracy'])
history = logreg_cls.fit(x_train.reshape(n_train, -1),
                        y_train,
                        batch_size = 512,
                        validation_data=[x_val.reshape(n_val, -1), y_val],
                        validation_batch_size=128,
                        epochs = 100,
                        callbacks = [tf.keras.callbacks.EarlyStopping(monitor_
    ↳ 'val_loss', patience = 10),
                                tf.keras.callbacks.ModelCheckpoint(str =_
    ↳ 'val_loss', save_best_only = True, save_weights_only = True, filepath=_
    ↳ "LR_No_Aug")],
                        verbose = 1)

logreg_cls.load_weights('LR_No_Aug')
preds_train = np.argmax(logreg_cls.predict(x_train.reshape(n_train, -1)), axis_
    ↳ 1)
preds_val    = np.argmax(logreg_cls.predict(x_val.reshape(n_val, -1)), axis = 1)
preds_test   = np.argmax(logreg_cls.predict(x_test.reshape(n_test, -1)), axis =_
    ↳ 1)

fig, ax = plot_history(history)

fig.savefig(PLOT_SAVE + "logreg_unaug_losscurve.png", facecolor = 'white')

print("Training Classification Report: \n", classification_report(y_train,_
    ↳ preds_train, digits = N_DIGITS))
print("\nValidation Classification Report: \n", classification_report(y_val,_
    ↳ preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test,_
    ↳ preds_test, digits = N_DIGITS))
```

Epoch 1/100



41/41 [=====] - 1s 6ms/step - loss: 0.7001 - accuracy:  
0.5477 - val\_loss: 0.6641 - val\_accuracy: 0.5994  
Epoch 2/100  
41/41 [=====] - 0s 3ms/step - loss: 0.6591 - accuracy:  
0.6098 - val\_loss: 0.6541 - val\_accuracy: 0.6092  
Epoch 3/100  
41/41 [=====] - 0s 3ms/step - loss: 0.6455 - accuracy:  
0.6239 - val\_loss: 0.6298 - val\_accuracy: 0.6453  
Epoch 4/100  
41/41 [=====] - 0s 3ms/step - loss: 0.6295 - accuracy:  
0.6464 - val\_loss: 0.6219 - val\_accuracy: 0.6510  
Epoch 5/100  
41/41 [=====] - 0s 3ms/step - loss: 0.6230 - accuracy:  
0.6522 - val\_loss: 0.6134 - val\_accuracy: 0.6643  
Epoch 6/100  
41/41 [=====] - 0s 3ms/step - loss: 0.6151 - accuracy:  
0.6637 - val\_loss: 0.6096 - val\_accuracy: 0.6714  
Epoch 7/100  
41/41 [=====] - 0s 3ms/step - loss: 0.6121 - accuracy:  
0.6664 - val\_loss: 0.6055 - val\_accuracy: 0.6724  
Epoch 8/100  
41/41 [=====] - 0s 3ms/step - loss: 0.6087 - accuracy:  
0.6705 - val\_loss: 0.6036 - val\_accuracy: 0.6772  
Epoch 9/100  
41/41 [=====] - 0s 3ms/step - loss: 0.6068 - accuracy:  
0.6682 - val\_loss: 0.6149 - val\_accuracy: 0.6612  
Epoch 10/100  
41/41 [=====] - 0s 3ms/step - loss: 0.6092 - accuracy:  
0.6697 - val\_loss: 0.6094 - val\_accuracy: 0.6704  
Epoch 11/100  
41/41 [=====] - 0s 4ms/step - loss: 0.6058 - accuracy:  
0.6729 - val\_loss: 0.6008 - val\_accuracy: 0.6796  
Epoch 12/100  
41/41 [=====] - 0s 3ms/step - loss: 0.6095 - accuracy:  
0.6693 - val\_loss: 0.6073 - val\_accuracy: 0.6779  
Epoch 13/100  
41/41 [=====] - 0s 3ms/step - loss: 0.6148 - accuracy:  
0.6622 - val\_loss: 0.6090 - val\_accuracy: 0.6741  
Epoch 14/100  
41/41 [=====] - 0s 3ms/step - loss: 0.6088 - accuracy:  
0.6721 - val\_loss: 0.6100 - val\_accuracy: 0.6718  
Epoch 15/100  
41/41 [=====] - 0s 3ms/step - loss: 0.6029 - accuracy:  
0.6750 - val\_loss: 0.5984 - val\_accuracy: 0.6837  
Epoch 16/100  
41/41 [=====] - 0s 3ms/step - loss: 0.6008 - accuracy:  
0.6795 - val\_loss: 0.5996 - val\_accuracy: 0.6837  
Epoch 17/100

41/41 [=====] - 0s 3ms/step - loss: 0.6022 - accuracy: 0.6776 - val\_loss: 0.5996 - val\_accuracy: 0.6809  
Epoch 18/100  
41/41 [=====] - 0s 4ms/step - loss: 0.6005 - accuracy: 0.6813 - val\_loss: 0.5970 - val\_accuracy: 0.6867  
Epoch 19/100  
41/41 [=====] - 0s 3ms/step - loss: 0.6010 - accuracy: 0.6802 - val\_loss: 0.5992 - val\_accuracy: 0.6867  
Epoch 20/100  
41/41 [=====] - 0s 3ms/step - loss: 0.6014 - accuracy: 0.6787 - val\_loss: 0.6112 - val\_accuracy: 0.6680  
Epoch 21/100  
41/41 [=====] - 0s 3ms/step - loss: 0.6045 - accuracy: 0.6763 - val\_loss: 0.6003 - val\_accuracy: 0.6840  
Epoch 22/100  
41/41 [=====] - 0s 3ms/step - loss: 0.6027 - accuracy: 0.6758 - val\_loss: 0.6090 - val\_accuracy: 0.6741  
Epoch 23/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5996 - accuracy: 0.6813 - val\_loss: 0.5960 - val\_accuracy: 0.6864  
Epoch 24/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5981 - accuracy: 0.6836 - val\_loss: 0.5957 - val\_accuracy: 0.6854  
Epoch 25/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5979 - accuracy: 0.6823 - val\_loss: 0.5986 - val\_accuracy: 0.6854  
Epoch 26/100  
41/41 [=====] - 0s 4ms/step - loss: 0.6000 - accuracy: 0.6783 - val\_loss: 0.5996 - val\_accuracy: 0.6843  
Epoch 27/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5971 - accuracy: 0.6835 - val\_loss: 0.5966 - val\_accuracy: 0.6860  
Epoch 28/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5963 - accuracy: 0.6844 - val\_loss: 0.5948 - val\_accuracy: 0.6888  
Epoch 29/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5949 - accuracy: 0.6841 - val\_loss: 0.5962 - val\_accuracy: 0.6915  
Epoch 30/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5948 - accuracy: 0.6857 - val\_loss: 0.6099 - val\_accuracy: 0.6731  
Epoch 31/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5974 - accuracy: 0.6808 - val\_loss: 0.5952 - val\_accuracy: 0.6894  
Epoch 32/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5933 - accuracy: 0.6882 - val\_loss: 0.6008 - val\_accuracy: 0.6782  
Epoch 33/100


41/41 [=====] - 0s 3ms/step - loss: 0.5961 - accuracy:  
0.6845 - val\_loss: 0.5962 - val\_accuracy: 0.6871  
Epoch 34/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5950 - accuracy:  
0.6863 - val\_loss: 0.5970 - val\_accuracy: 0.6881  
Epoch 35/100  
41/41 [=====] - 0s 3ms/step - loss: 0.6029 - accuracy:  
0.6764 - val\_loss: 0.5937 - val\_accuracy: 0.6898  
Epoch 36/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5977 - accuracy:  
0.6845 - val\_loss: 0.6002 - val\_accuracy: 0.6843  
Epoch 37/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5965 - accuracy:  
0.6857 - val\_loss: 0.5997 - val\_accuracy: 0.6847  
Epoch 38/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5945 - accuracy:  
0.6834 - val\_loss: 0.6005 - val\_accuracy: 0.6826  
Epoch 39/100  
41/41 [=====] - 0s 4ms/step - loss: 0.5930 - accuracy:  
0.6857 - val\_loss: 0.5954 - val\_accuracy: 0.6908  
Epoch 40/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5942 - accuracy:  
0.6850 - val\_loss: 0.5983 - val\_accuracy: 0.6830  
Epoch 41/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5938 - accuracy:  
0.6872 - val\_loss: 0.5956 - val\_accuracy: 0.6854  
Epoch 42/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5914 - accuracy:  
0.6876 - val\_loss: 0.6017 - val\_accuracy: 0.6809  
Epoch 43/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5937 - accuracy:  
0.6861 - val\_loss: 0.5976 - val\_accuracy: 0.6854  
Epoch 44/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5912 - accuracy:  
0.6877 - val\_loss: 0.5966 - val\_accuracy: 0.6894  
Epoch 45/100  
41/41 [=====] - 0s 4ms/step - loss: 0.5901 - accuracy:  
0.6902 - val\_loss: 0.5932 - val\_accuracy: 0.6905  
Epoch 46/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5895 - accuracy:  
0.6908 - val\_loss: 0.5983 - val\_accuracy: 0.6826  
Epoch 47/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5963 - accuracy:  
0.6841 - val\_loss: 0.6039 - val\_accuracy: 0.6775  
Epoch 48/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5923 - accuracy:  
0.6889 - val\_loss: 0.5942 - val\_accuracy: 0.6884  
Epoch 49/100

41/41 [=====] - 0s 4ms/step - loss: 0.5914 - accuracy:  
0.6906 - val\_loss: 0.5931 - val\_accuracy: 0.6874  
Epoch 50/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5881 - accuracy:  
0.6942 - val\_loss: 0.5912 - val\_accuracy: 0.6925  
Epoch 51/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5893 - accuracy:  
0.6891 - val\_loss: 0.5969 - val\_accuracy: 0.6857  
Epoch 52/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5912 - accuracy:  
0.6886 - val\_loss: 0.5925 - val\_accuracy: 0.6888  
Epoch 53/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5900 - accuracy:  
0.6891 - val\_loss: 0.5928 - val\_accuracy: 0.6898  
Epoch 54/100  
41/41 [=====] - 0s 4ms/step - loss: 0.5901 - accuracy:  
0.6916 - val\_loss: 0.5944 - val\_accuracy: 0.6837  
Epoch 55/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5926 - accuracy:  
0.6889 - val\_loss: 0.5961 - val\_accuracy: 0.6843  
Epoch 56/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5896 - accuracy:  
0.6893 - val\_loss: 0.5919 - val\_accuracy: 0.6925  
Epoch 57/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5908 - accuracy:  
0.6884 - val\_loss: 0.5911 - val\_accuracy: 0.6881  
Epoch 58/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5936 - accuracy:  
0.6873 - val\_loss: 0.6000 - val\_accuracy: 0.6806  
Epoch 59/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5878 - accuracy:  
0.6936 - val\_loss: 0.5903 - val\_accuracy: 0.6955  
Epoch 60/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5880 - accuracy:  
0.6906 - val\_loss: 0.5916 - val\_accuracy: 0.6955  
Epoch 61/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5912 - accuracy:  
0.6894 - val\_loss: 0.5895 - val\_accuracy: 0.6959  
Epoch 62/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5905 - accuracy:  
0.6910 - val\_loss: 0.5900 - val\_accuracy: 0.6928  
Epoch 63/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5859 - accuracy:  
0.6950 - val\_loss: 0.5950 - val\_accuracy: 0.6854  
Epoch 64/100  
41/41 [=====] - 0s 3ms/step - loss: 0.5880 - accuracy:  
0.6911 - val\_loss: 0.5920 - val\_accuracy: 0.6938  
Epoch 65/100

```

41/41 [=====] - 0s 3ms/step - loss: 0.5918 - accuracy:
0.6875 - val_loss: 0.6114 - val_accuracy: 0.6762
Epoch 66/100
41/41 [=====] - 0s 3ms/step - loss: 0.5877 - accuracy:
0.6920 - val_loss: 0.5995 - val_accuracy: 0.6789
Epoch 67/100
41/41 [=====] - 0s 3ms/step - loss: 0.5861 - accuracy:
0.6929 - val_loss: 0.5981 - val_accuracy: 0.6820
Epoch 68/100
41/41 [=====] - 0s 3ms/step - loss: 0.5888 - accuracy:
0.6881 - val_loss: 0.5908 - val_accuracy: 0.6972
Epoch 69/100
41/41 [=====] - 0s 4ms/step - loss: 0.5854 - accuracy:
0.6960 - val_loss: 0.6110 - val_accuracy: 0.6639
Epoch 70/100
41/41 [=====] - 0s 4ms/step - loss: 0.5920 - accuracy:
0.6868 - val_loss: 0.6058 - val_accuracy: 0.6779
Epoch 71/100
41/41 [=====] - 0s 3ms/step - loss: 0.5883 - accuracy:
0.6896 - val_loss: 0.5957 - val_accuracy: 0.6864
644/644 [=====] - 0s 701us/step
92/92 [=====] - 0s 709us/step
82/82 [=====] - 0s 716us/step

```



Training Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0.0	0.701	0.687	0.694	10260
1.0	0.695	0.710	0.703	10341
accuracy			0.698	20601
macro avg	0.698	0.698	0.698	20601
weighted avg	0.698	0.698	0.698	20601

#### Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.696	0.692	0.694	1466
1.0	0.696	0.699	0.698	1477
accuracy			0.696	2943
macro avg	0.696	0.696	0.696	2943
weighted avg	0.696	0.696	0.696	2943

#### Testing Classification Report:

	precision	recall	f1-score	support
0.0	0.693	0.691	0.692	1303
1.0	0.695	0.696	0.696	1314
accuracy			0.694	2617
macro avg	0.694	0.694	0.694	2617
weighted avg	0.694	0.694	0.694	2617

### Decision Trees

```
[ ]: from sklearn.tree import DecisionTreeClassifier

dt_cls = DecisionTreeClassifier(criterion = 'gini', splitter = 'best',
    ↳max_depth = 5, min_samples_split = 2, )
dt_cls.fit(x_train.reshape(n_train, -1), y_train)
preds_train = dt_cls.predict(x_train.reshape(n_train, -1))
preds_val   = dt_cls.predict(x_val.reshape(n_val, -1))
preds_test  = dt_cls.predict(x_test.reshape(n_test, -1))

print("Training Classification Report: \n", classification_report(y_train,
    ↳preds_train, digits = N_DIGITS))
print("\nValidation Classification Report: \n", classification_report(y_val,
    ↳preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test,
    ↳preds_test, digits = N_DIGITS))
```

```
# eval = testing_module.ModelEvaluation(y_val,
#     nb_cls.predict(x_val.reshape(n_val, -1)),
#     model_reference_name = 'Naive Bayes',
#     model_type = 'classification',
#     plot_classification_metric = ['roc_auc']) # if classification
# eval.evaluate(evaluate_save= True, plots_show = True)
```

#### Training Classification Report:

	precision	recall	f1-score	support
0.0	0.685	0.774	0.727	10260
1.0	0.743	0.646	0.691	10341
accuracy			0.710	20601
macro avg	0.714	0.710	0.709	20601
weighted avg	0.714	0.710	0.709	20601

#### Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.676	0.774	0.721	1466
1.0	0.738	0.632	0.681	1477
accuracy			0.702	2943
macro avg	0.707	0.703	0.701	2943
weighted avg	0.707	0.702	0.701	2943

#### Testing Classification Report:

	precision	recall	f1-score	support
0.0	0.661	0.748	0.702	1303
1.0	0.713	0.619	0.663	1314
accuracy			0.684	2617
macro avg	0.687	0.684	0.682	2617
weighted avg	0.687	0.684	0.682	2617

### XGBoost

```
[ ]: tf.keras.backend.clear_session()

from xgboost import XGBClassifier

xgb_cls = XGBClassifier(max_depth = 5, objective = 'reg:logistic',
                        num_parallel_tree = 20, booster = 'gbtree',
```

```

gamma = 0.5, tree_method = 'gpu_hist', subsample = 0.4, reg_lambda = 1)
xgb_cls.fit(x_train.reshape(n_train, -1), y_train)
preds_train = xgb_cls.predict(x_train.reshape(n_train, -1))
preds_val = xgb_cls.predict(x_val.reshape(n_val, -1))
preds_test = xgb_cls.predict(x_test.reshape(n_test, -1))

print("Training Classification Report: \n", classification_report(y_train,
    preds_train, digits = N_DIGITS))
print("\nValidation Classification Report: \n", classification_report(y_val,
    preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test,
    preds_test, digits = N_DIGITS))

```

#### Training Classification Report:

	precision	recall	f1-score	support
0.0	0.966	0.983	0.974	10260
1.0	0.983	0.966	0.974	10341
accuracy			0.974	20601
macro avg	0.974	0.974	0.974	20601
weighted avg	0.975	0.974	0.974	20601

#### Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.842	0.873	0.857	1466
1.0	0.869	0.837	0.853	1477
accuracy			0.855	2943
macro avg	0.855	0.855	0.855	2943
weighted avg	0.855	0.855	0.855	2943

#### Testing Classification Report:

	precision	recall	f1-score	support
0.0	0.842	0.861	0.851	1303
1.0	0.859	0.839	0.849	1314
accuracy			0.850	2617
macro avg	0.850	0.850	0.850	2617
weighted avg	0.850	0.850	0.850	2617



## SVM

```
[ ]: from sklearn.svm import SVC

svm_cls = SVC(kernel = 'poly', degree = 3, gamma = 'auto', max_iter = 250,
↳verbose= True)
svm_cls.fit(x_train.reshape(n_train, -1), y_train)
preds_train = svm_cls.predict(x_train.reshape(n_train, -1))
preds_val   = svm_cls.predict(x_val.reshape(n_val, -1))
preds_test  = svm_cls.predict(x_test.reshape(n_test, -1))

print("Training Classification Report: \n", classification_report(y_train,
↳preds_train, digits = N_DIGITS))
print("\nValidation Classification Report: \n", classification_report(y_val,
↳preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test,
↳preds_test, digits = N_DIGITS))
```

[LibSVM]

```
c:\ProgramData\Anaconda3\envs\ml\lib\site-packages\sklearn\svm\_base.py:301:
ConvergenceWarning: Solver terminated early (max_iter=250). Consider pre-
processing your data with StandardScaler or MinMaxScaler.
  warnings.warn(
```

Training Classification Report:

	precision	recall	f1-score	support
0.0	0.581	0.102	0.174	10260
1.0	0.510	0.927	0.658	10341
accuracy			0.516	20601
macro avg	0.545	0.515	0.416	20601
weighted avg	0.545	0.516	0.417	20601

Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.585	0.101	0.172	1466
1.0	0.510	0.929	0.659	1477
accuracy			0.516	2943
macro avg	0.548	0.515	0.415	2943
weighted avg	0.547	0.516	0.416	2943

Testing Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

	0.0	0.590	0.106	0.180	1303
	1.0	0.511	0.927	0.659	1314
accuracy				0.518	2617
macro avg	0.550	0.516	0.419		2617
weighted avg	0.550	0.518	0.420		2617

## Transfer Learning

```
[ ]: imagenet = tf.keras.applications.Xception(
    weights = 'imagenet',
    include_top = False,
    input_shape = (72, 72, 3)
)
imagenet.trainable = False

trans_learn = tf.keras.Sequential([
    tf.keras.layers.Resizing(72, 72),
    imagenet,

    tf.keras.layers.GlobalAveragePooling2D(),
    tf.keras.layers.Dense(512, activation = 'relu'),
    tf.keras.layers.Dense(512, activation = 'relu'),
    tf.keras.layers.Dense(2, activation = 'sigmoid')
])

trans_learn.compile(
    optimizer = tf.keras.optimizers.Adam(learning_rate=0.003),
    loss = 'sparse_categorical_crossentropy',
    metrics = ['accuracy']
)

history = trans_learn.fit(
    x_train, y_train, batch_size = 64,
    shuffle = True,
    epochs = 50,
    validation_data = [x_val, y_val],
    validation_batch_size = 32,
    callbacks = [tf.keras.callbacks.EarlyStopping(monitor = 'val_loss',
    ↪patience = 10),
                 tf.keras.callbacks.ModelCheckpoint(str = 'val_loss',
    ↪save_best_only = True, save_weights_only = True, filepath= "TL_No_Aug")]
    )
fig, ax = plot_history(history)

trans_learn.load_weights('TL_No_Aug')
```

```

fig.savefig(PLOT_SAVE + "TL_unaug_losscurve.png", facecolor = 'white')
preds_train = np.argmax(trans_learn.predict(x_train), axis = 1)
preds_val    = np.argmax(trans_learn.predict(x_val), axis = 1)
preds_test   = np.argmax(trans_learn.predict(x_test), axis = 1)

print("Training Classification Report: \n", classification_report(y_train,
↳preds_train, digits = N_DIGITS))
print("\nValidation Classification Report: \n", classification_report(y_val,
↳preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test,
↳preds_test, digits = N_DIGITS))

```

Epoch 1/50

322/322 [=====] - 9s 19ms/step - loss: 0.2352 -  
accuracy: 0.9161 - val\_loss: 0.1279 - val\_accuracy: 0.9531

Epoch 2/50

322/322 [=====] - 5s 16ms/step - loss: 0.1295 -  
accuracy: 0.9490 - val\_loss: 0.1456 - val\_accuracy: 0.9439

Epoch 3/50

322/322 [=====] - 5s 16ms/step - loss: 0.1155 -  
accuracy: 0.9555 - val\_loss: 0.1284 - val\_accuracy: 0.9501

Epoch 4/50

322/322 [=====] - 5s 16ms/step - loss: 0.1016 -  
accuracy: 0.9597 - val\_loss: 0.1136 - val\_accuracy: 0.9616

Epoch 5/50

322/322 [=====] - 5s 16ms/step - loss: 0.0926 -  
accuracy: 0.9641 - val\_loss: 0.1082 - val\_accuracy: 0.9626

Epoch 6/50

322/322 [=====] - 5s 15ms/step - loss: 0.0841 -  
accuracy: 0.9680 - val\_loss: 0.1310 - val\_accuracy: 0.9609

Epoch 7/50

322/322 [=====] - 5s 15ms/step - loss: 0.0837 -  
accuracy: 0.9676 - val\_loss: 0.1182 - val\_accuracy: 0.9558

Epoch 8/50

322/322 [=====] - 5s 15ms/step - loss: 0.0769 -  
accuracy: 0.9707 - val\_loss: 0.1169 - val\_accuracy: 0.9650

Epoch 9/50

322/322 [=====] - 5s 15ms/step - loss: 0.0694 -  
accuracy: 0.9728 - val\_loss: 0.1305 - val\_accuracy: 0.9575

Epoch 10/50

322/322 [=====] - 5s 15ms/step - loss: 0.0642 -  
accuracy: 0.9754 - val\_loss: 0.1093 - val\_accuracy: 0.9613

Epoch 11/50


322/322 [=====] - 5s 15ms/step - loss: 0.0622 -  
accuracy: 0.9757 - val\_loss: 0.1323 - val\_accuracy: 0.9647

Epoch 12/50

```

322/322 [=====] - 5s 15ms/step - loss: 0.0587 -
accuracy: 0.9767 - val_loss: 0.1749 - val_accuracy: 0.9494
Epoch 13/50
322/322 [=====] - 5s 15ms/step - loss: 0.0586 -
accuracy: 0.9784 - val_loss: 0.1238 - val_accuracy: 0.9613
Epoch 14/50
322/322 [=====] - 5s 15ms/step - loss: 0.0519 -
accuracy: 0.9804 - val_loss: 0.1277 - val_accuracy: 0.9613
Epoch 15/50
322/322 [=====] - 5s 15ms/step - loss: 0.0487 -
accuracy: 0.9819 - val_loss: 0.1478 - val_accuracy: 0.9541

```



```

644/644 [=====] - 5s 7ms/step
92/92 [=====] - 1s 7ms/step
82/82 [=====] - 1s 7ms/step

```

Training Classification Report:

	precision	recall	f1-score	support
0.0	0.959	0.985	0.972	10260
1.0	0.985	0.958	0.971	10341
accuracy			0.972	20601
macro avg	0.972	0.972	0.972	20601
weighted avg	0.972	0.972	0.972	20601

Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.952	0.974	0.963	1466
1.0	0.974	0.951	0.962	1477
accuracy			0.963	2943
macro avg	0.963	0.963	0.963	2943
weighted avg	0.963	0.963	0.963	2943

Testing Classification Report:

	precision	recall	f1-score	support
0.0	0.944	0.968	0.956	1303
1.0	0.967	0.943	0.955	1314
accuracy			0.955	2617
macro avg	0.956	0.955	0.955	2617
weighted avg	0.956	0.955	0.955	2617

## CNN Model

```
[ ]: cnn = tf.keras.Sequential([
    tf.keras.layers.Conv2D(16, (3,3), padding = 'same', activation = 'relu',
    ↪input_shape = x_train[0].shape),
    tf.keras.layers.MaxPool2D((3,3), padding = 'same'),

    tf.keras.layers.Conv2D(32, (2,2), padding = 'same', activation = 'relu'),
    tf.keras.layers.MaxPool2D((2,2), padding = 'same'),

    tf.keras.layers.Flatten(),

    tf.keras.layers.Dense(256, activation = 'relu'),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(256, activation = 'relu'),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(2, activation = 'sigmoid')
])

cnn.compile(
    optimizer = tf.keras.optimizers.Adam(learning_rate=0.003),
    loss = 'sparse_categorical_crossentropy',
    metrics = ['accuracy']
)

history = cnn.fit(
```

```

        x_train, y_train, batch_size = 1024,
        epochs = 50,
        validation_data = [x_val, y_val],
        validation_batch_size = 256,
        callbacks = [tf.keras.callbacks.EarlyStopping(monitor = 'val_loss',
↪patience = 10),
                    tf.keras.callbacks.ModelCheckpoint(str = 'val_loss',
↪save_best_only = True, save_weights_only = True, filepath= "CNN_No_Aug")]
    )

cnn.load_weights('CNN_No_Aug')
fig, ax = plot_history(history)

fig.savefig(PLOT_SAVE + "CNN_unaug_losscurve.png", facecolor = 'white')
preds_train = np.argmax(cnn.predict(x_train), axis = 1)
preds_val    = np.argmax(cnn.predict(x_val), axis = 1)
preds_test   = np.argmax(cnn.predict(x_test), axis = 1)

print("Training Classification Report: \n", classification_report(y_train,
↪preds_train, digits = N_DIGITS))
print("\nValidation Classification Report: \n", classification_report(y_val,
↪preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test,
↪preds_test, digits = N_DIGITS))

```

Epoch 1/50

21/21 [=====] - 1s 18ms/step - loss: 0.6625 - accuracy: 0.5956 - val\_loss: 0.6094 - val\_accuracy: 0.6758

Epoch 2/50

21/21 [=====] - 0s 9ms/step - loss: 0.6057 - accuracy: 0.6752 - val\_loss: 0.5299 - val\_accuracy: 0.7601

Epoch 3/50

21/21 [=====] - 0s 9ms/step - loss: 0.5441 - accuracy: 0.7230 - val\_loss: 0.5027 - val\_accuracy: 0.7666

Epoch 4/50

21/21 [=====] - 0s 9ms/step - loss: 0.4931 - accuracy: 0.7666 - val\_loss: 0.4683 - val\_accuracy: 0.7883

Epoch 5/50

21/21 [=====] - 0s 9ms/step - loss: 0.4335 - accuracy: 0.7987 - val\_loss: 0.3831 - val\_accuracy: 0.8264

Epoch 6/50

21/21 [=====] - 0s 7ms/step - loss: 0.4756 - accuracy: 0.7706 - val\_loss: 0.4782 - val\_accuracy: 0.7676

Epoch 7/50

21/21 [=====] - 0s 9ms/step - loss: 0.4180 - accuracy: 0.8054 - val\_loss: 0.3376 - val\_accuracy: 0.8488

Epoch 8/50  
21/21 [=====] - 0s 8ms/step - loss: 0.3718 - accuracy:  
0.8270 - val\_loss: 0.4914 - val\_accuracy: 0.7319  
Epoch 9/50  
21/21 [=====] - 0s 9ms/step - loss: 0.3703 - accuracy:  
0.8299 - val\_loss: 0.2982 - val\_accuracy: 0.8668  
Epoch 10/50  
21/21 [=====] - 0s 9ms/step - loss: 0.3300 - accuracy:  
0.8563 - val\_loss: 0.2942 - val\_accuracy: 0.8763  
Epoch 11/50  
21/21 [=====] - 0s 9ms/step - loss: 0.2771 - accuracy:  
0.8814 - val\_loss: 0.2473 - val\_accuracy: 0.8943  
Epoch 12/50  
21/21 [=====] - 0s 9ms/step - loss: 0.2127 - accuracy:  
0.9124 - val\_loss: 0.2012 - val\_accuracy: 0.9188  
Epoch 13/50  
21/21 [=====] - 0s 7ms/step - loss: 0.1878 - accuracy:  
0.9268 - val\_loss: 0.2146 - val\_accuracy: 0.9171  
Epoch 14/50  
21/21 [=====] - 0s 9ms/step - loss: 0.1714 - accuracy:  
0.9360 - val\_loss: 0.1777 - val\_accuracy: 0.9348  
Epoch 15/50  
21/21 [=====] - 0s 9ms/step - loss: 0.1311 - accuracy:  
0.9503 - val\_loss: 0.1491 - val\_accuracy: 0.9422  
Epoch 16/50  
21/21 [=====] - 0s 9ms/step - loss: 0.1288 - accuracy:  
0.9523 - val\_loss: 0.1300 - val\_accuracy: 0.9504  
Epoch 17/50  
21/21 [=====] - 0s 9ms/step - loss: 0.1055 - accuracy:  
0.9607 - val\_loss: 0.0972 - val\_accuracy: 0.9684  
Epoch 18/50  
21/21 [=====] - 0s 9ms/step - loss: 0.0907 - accuracy:  
0.9663 - val\_loss: 0.0963 - val\_accuracy: 0.9687  
Epoch 19/50  
21/21 [=====] - 0s 9ms/step - loss: 0.0802 - accuracy:  
0.9711 - val\_loss: 0.0801 - val\_accuracy: 0.9728  
Epoch 20/50  
21/21 [=====] - 0s 8ms/step - loss: 0.0728 - accuracy:  
0.9731 - val\_loss: 0.0790 - val\_accuracy: 0.9718  
Epoch 21/50  
21/21 [=====] - 0s 7ms/step - loss: 0.0634 - accuracy:  
0.9775 - val\_loss: 0.0807 - val\_accuracy: 0.9704  
Epoch 22/50  
21/21 [=====] - 0s 8ms/step - loss: 0.0642 - accuracy:  
0.9759 - val\_loss: 0.0795 - val\_accuracy: 0.9711  
Epoch 23/50  
21/21 [=====] - 0s 9ms/step - loss: 0.0511 - accuracy:  
0.9817 - val\_loss: 0.0736 - val\_accuracy: 0.9742

Epoch 24/50  
21/21 [=====] - 0s 8ms/step - loss: 0.0516 - accuracy: 0.9819 - val\_loss: 0.0839 - val\_accuracy: 0.9732

Epoch 25/50  
21/21 [=====] - 0s 9ms/step - loss: 0.0528 - accuracy: 0.9820 - val\_loss: 0.0708 - val\_accuracy: 0.9783

Epoch 26/50  
21/21 [=====] - 0s 9ms/step - loss: 0.0511 - accuracy: 0.9833 - val\_loss: 0.0681 - val\_accuracy: 0.9755

Epoch 27/50  
21/21 [=====] - 0s 7ms/step - loss: 0.0536 - accuracy: 0.9811 - val\_loss: 0.0696 - val\_accuracy: 0.9769

Epoch 28/50  
21/21 [=====] - 0s 9ms/step - loss: 0.0426 - accuracy: 0.9850 - val\_loss: 0.0609 - val\_accuracy: 0.9820

Epoch 29/50  
21/21 [=====] - 0s 9ms/step - loss: 0.0393 - accuracy: 0.9861 - val\_loss: 0.0572 - val\_accuracy: 0.9813

Epoch 30/50  
21/21 [=====] - 0s 7ms/step - loss: 0.0378 - accuracy: 0.9869 - val\_loss: 0.0635 - val\_accuracy: 0.9803

Epoch 31/50  
21/21 [=====] - 0s 9ms/step - loss: 0.0447 - accuracy: 0.9842 - val\_loss: 0.0560 - val\_accuracy: 0.9817

Epoch 32/50  
21/21 [=====] - 0s 9ms/step - loss: 0.0326 - accuracy: 0.9890 - val\_loss: 0.0529 - val\_accuracy: 0.9834

Epoch 33/50  
21/21 [=====] - 0s 9ms/step - loss: 0.0291 - accuracy: 0.9901 - val\_loss: 0.0522 - val\_accuracy: 0.9837

Epoch 34/50  
21/21 [=====] - 0s 9ms/step - loss: 0.0228 - accuracy: 0.9924 - val\_loss: 0.0490 - val\_accuracy: 0.9857

Epoch 35/50  
21/21 [=====] - 0s 8ms/step - loss: 0.0216 - accuracy: 0.9931 - val\_loss: 0.0548 - val\_accuracy: 0.9857

Epoch 36/50  
21/21 [=====] - 0s 7ms/step - loss: 0.0252 - accuracy: 0.9914 - val\_loss: 0.0538 - val\_accuracy: 0.9840


Epoch 37/50  
21/21 [=====] - 0s 9ms/step - loss: 0.0245 - accuracy: 0.9908 - val\_loss: 0.0459 - val\_accuracy: 0.9857

Epoch 38/50  
21/21 [=====] - 0s 7ms/step - loss: 0.0165 - accuracy: 0.9945 - val\_loss: 0.0561 - val\_accuracy: 0.9847

Epoch 39/50  
21/21 [=====] - 0s 7ms/step - loss: 0.0219 - accuracy: 0.9919 - val\_loss: 0.0651 - val\_accuracy: 0.9834



Epoch 40/50  
 21/21 [=====] - 0s 8ms/step - loss: 0.0213 - accuracy: 0.9931 - val\_loss: 0.0574 - val\_accuracy: 0.9854  
 Epoch 41/50  
 21/21 [=====] - 0s 8ms/step - loss: 0.0220 - accuracy: 0.9919 - val\_loss: 0.0588 - val\_accuracy: 0.9847  
 Epoch 42/50  
 21/21 [=====] - 0s 7ms/step - loss: 0.0277 - accuracy: 0.9902 - val\_loss: 0.0579 - val\_accuracy: 0.9861  
 Epoch 43/50  
 21/21 [=====] - 0s 8ms/step - loss: 0.0217 - accuracy: 0.9930 - val\_loss: 0.0520 - val\_accuracy: 0.9847  
 Epoch 44/50  
 21/21 [=====] - 0s 7ms/step - loss: 0.0197 - accuracy: 0.9936 - val\_loss: 0.0501 - val\_accuracy: 0.9847  
 Epoch 45/50  
 21/21 [=====] - 0s 7ms/step - loss: 0.0147 - accuracy: 0.9953 - val\_loss: 0.0702 - val\_accuracy: 0.9837  
 Epoch 46/50  
 21/21 [=====] - 0s 7ms/step - loss: 0.0146 - accuracy: 0.9953 - val\_loss: 0.0640 - val\_accuracy: 0.9850  
 Epoch 47/50  
 21/21 [=====] - 0s 8ms/step - loss: 0.0173 - accuracy: 0.9941 - val\_loss: 0.0785 - val\_accuracy: 0.9827



644/644 [=====] - 1s 1ms/step  
 92/92 [=====] - 0s 1ms/step

82/82 [=====] - 0s 1ms/step

#### Training Classification Report:

	precision	recall	f1-score	support
0.0	0.994	0.997	0.996	10260
1.0	0.997	0.994	0.996	10341
accuracy			0.996	20601
macro avg	0.996	0.996	0.996	20601
weighted avg	0.996	0.996	0.996	20601

#### Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.983	0.988	0.986	1466
1.0	0.988	0.983	0.986	1477
accuracy			0.986	2943
macro avg	0.986	0.986	0.986	2943
weighted avg	0.986	0.986	0.986	2943

#### Testing Classification Report:

	precision	recall	f1-score	support
0.0	0.976	0.988	0.982	1303
1.0	0.988	0.976	0.982	1314
accuracy			0.982	2617
macro avg	0.982	0.982	0.982	2617
weighted avg	0.982	0.982	0.982	2617

## 0.6.2 Augmented Dataset

### Naive Bayes

```
[ ]: from sklearn import naive_bayes

nb_cls = naive_bayes.GaussianNB()
nb_cls.fit(x_train_aug.reshape(n_aug_train, -1), y_train_aug)
preds_train = nb_cls.predict(x_train_aug.reshape(n_aug_train, -1))
preds_val   = nb_cls.predict(x_val.reshape(n_val, -1))
preds_test  = nb_cls.predict(x_test.reshape(n_test, -1))

print("Training Classification Report: \n", classification_report(y_train_aug,
↪preds_train, digits = N_DIGITS))
```

```
print("\nValidation Classification Report: \n", classification_report(y_val,
↪preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test,
↪preds_test, digits = N_DIGITS))
```

#### Training Classification Report:

	precision	recall	f1-score	support
0.0	0.607	0.760	0.675	20520
1.0	0.683	0.512	0.586	20682
accuracy			0.636	41202
macro avg	0.645	0.636	0.631	41202
weighted avg	0.645	0.636	0.630	41202

#### Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.000	0.000	0.000	1466
1.0	0.502	1.000	0.668	1477
accuracy			0.502	2943
macro avg	0.251	0.500	0.334	2943
weighted avg	0.252	0.502	0.335	2943

#### Testing Classification Report:

	precision	recall	f1-score	support
0.0	0.000	0.000	0.000	1303
1.0	0.502	1.000	0.669	1314
accuracy			0.502	2617
macro avg	0.251	0.500	0.334	2617
weighted avg	0.252	0.502	0.336	2617

```
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
```

```

c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))

```

## Logistic Regression

```

[ ]: from sklearn.linear_model import LogisticRegression

logreg_cls = tf.keras.Sequential([tf.keras.layers.Dense(2, activation =_
    ↪ 'sigmoid', input_dim = HEIGHT * WIDTH * 3)])
logreg_cls.compile(optimizer = tf.optimizers.Adam(0.001), loss =_
    ↪ 'sparse_categorical_crossentropy', metrics = ['accuracy'])
history = logreg_cls.fit(x_train_aug.reshape(n_aug_train, -1),
                        y_train_aug,
                        batch_size = 512,
                        validation_data=[x_val.reshape(n_val, -1), y_val],
                        validation_batch_size=128,
                        epochs = 100,
                        callbacks = [tf.keras.callbacks.EarlyStopping(monitor_
    ↪ 'val_loss', patience = 10),
                                tf.keras.callbacks.ModelCheckpoint(str =_
    ↪ 'val_loss', save_best_only = True, save_weights_only = True, filepath=_
    ↪ "LR_Aug")])
preds_train = np.argmax(logreg_cls.predict(x_train_aug.reshape(n_aug_train,_
    ↪ -1)), axis = 1)
preds_val   = np.argmax(logreg_cls.predict(x_val.reshape(n_val, -1)), axis = 1)
preds_test  = np.argmax(logreg_cls.predict(x_test.reshape(n_test, -1)), axis =_
    ↪ 1)

logreg_cls.load_weights('LR_Aug')
fig, ax = plot_history(history)

```

```

fig.savefig(PLOT_SAVE + "logreg_aug_losscurve.png", facecolor = 'white')

print("Training Classification Report: \n", classification_report(y_train_aug,
    ↪preds_train, digits = N_DIGITS))
print("\nValidation Classification Report: \n", classification_report(y_val,
    ↪preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test,
    ↪preds_test, digits = N_DIGITS))

```

Epoch 1/100

81/81 [=====] - 0s 3ms/step - loss: 27.6512 - accuracy:  
0.5376 - val\_loss: 0.6819 - val\_accuracy: 0.6007

Epoch 2/100

81/81 [=====] - 0s 2ms/step - loss: 10.8976 - accuracy:  
0.5726 - val\_loss: 0.6814 - val\_accuracy: 0.6055

Epoch 3/100

81/81 [=====] - 0s 2ms/step - loss: 21.7795 - accuracy:  
0.5545 - val\_loss: 0.6752 - val\_accuracy: 0.6466

Epoch 4/100

81/81 [=====] - 0s 2ms/step - loss: 8.6269 - accuracy:  
0.6051 - val\_loss: 0.6807 - val\_accuracy: 0.6432

Epoch 5/100

81/81 [=====] - 0s 2ms/step - loss: 8.6525 - accuracy:  
0.5836 - val\_loss: 0.6813 - val\_accuracy: 0.6075

Epoch 6/100

81/81 [=====] - 0s 3ms/step - loss: 11.2940 - accuracy:  
0.5716 - val\_loss: 0.6885 - val\_accuracy: 0.5022

Epoch 7/100

81/81 [=====] - 0s 2ms/step - loss: 11.3517 - accuracy:  
0.5891 - val\_loss: 0.6810 - val\_accuracy: 0.6018

Epoch 8/100

81/81 [=====] - 0s 2ms/step - loss: 8.9236 - accuracy:  
0.5885 - val\_loss: 0.7014 - val\_accuracy: 0.5012

Epoch 9/100

81/81 [=====] - 0s 2ms/step - loss: 16.4641 - accuracy:  
0.5702 - val\_loss: 0.6786 - val\_accuracy: 0.6041

Epoch 10/100

81/81 [=====] - 0s 2ms/step - loss: 7.6863 - accuracy:  
0.6055 - val\_loss: 0.6820 - val\_accuracy: 0.6256

Epoch 11/100

81/81 [=====] - 0s 2ms/step - loss: 9.2223 - accuracy:  
0.5833 - val\_loss: 0.6793 - val\_accuracy: 0.6555

Epoch 12/100

81/81 [=====] - 0s 2ms/step - loss: 17.7575 - accuracy:  
0.5690 - val\_loss: 0.6769 - val\_accuracy: 0.5790

Epoch 13/100

```

81/81 [=====] - 0s 3ms/step - loss: 11.1797 - accuracy:
0.5774 - val_loss: 0.6755 - val_accuracy: 0.6092
1288/1288 [=====] - 1s 712us/step
92/92 [=====] - 0s 769us/step
82/82 [=====] - 0s 784us/step

```



#### Training Classification Report:

	precision	recall	f1-score	support
0.0	0.638	0.430	0.514	20520
1.0	0.573	0.759	0.653	20682
accuracy			0.595	41202
macro avg	0.606	0.594	0.583	41202
weighted avg	0.605	0.595	0.583	41202

#### Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.813	0.280	0.416	1466
1.0	0.567	0.936	0.706	1477
accuracy			0.609	2943
macro avg	0.690	0.608	0.561	2943
weighted avg	0.690	0.609	0.562	2943

### Testing Classification Report:

	precision	recall	f1-score	support
0.0	0.776	0.243	0.370	1303
1.0	0.553	0.931	0.694	1314
accuracy			0.588	2617
macro avg	0.665	0.587	0.532	2617
weighted avg	0.664	0.588	0.533	2617

### Decision Trees

```
[ ]: from sklearn.tree import DecisionTreeClassifier

dt_cls = DecisionTreeClassifier(criterion = 'gini', splitter = 'best',
    ↳max_depth = 5, min_samples_split = 2, )
dt_cls.fit(x_train_aug.reshape(n_aug_train, -1), y_train_aug)
preds_train = dt_cls.predict(x_train_aug.reshape(n_aug_train, -1))
preds_val   = dt_cls.predict(x_val.reshape(n_val, -1))
preds_test  = dt_cls.predict(x_test.reshape(n_test, -1))

print("Training Classification Report: \n", classification_report(y_train_aug,
    ↳preds_train, digits = N_DIGITS))
print("\nValidation Classification Report: \n", classification_report(y_val,
    ↳preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test,
    ↳preds_test, digits = N_DIGITS))

# eval = testing_module.ModelEvaluation(y_val,
#     nb_cls.predict(x_val.reshape(n_val, -1)),
#     model_reference_name = 'Naive Bayes',
#     model_type = 'classification',
#     plot_classification_metric = ['roc_auc']) # if classification
# eval.evaluate(evaluate_save= True, plots_show = True)
```

### Training Classification Report:

	precision	recall	f1-score	support
0.0	0.687	0.768	0.725	20520
1.0	0.739	0.653	0.693	20682
accuracy			0.710	41202
macro avg	0.713	0.710	0.709	41202
weighted avg	0.713	0.710	0.709	41202

### Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.000	0.000	0.000	1466
1.0	0.502	1.000	0.668	1477
accuracy			0.502	2943
macro avg	0.251	0.500	0.334	2943
weighted avg	0.252	0.502	0.335	2943

#### Testing Classification Report:

	precision	recall	f1-score	support
0.0	0.000	0.000	0.000	1303
1.0	0.502	1.000	0.669	1314
accuracy			0.502	2617
macro avg	0.251	0.500	0.334	2617
weighted avg	0.252	0.502	0.336	2617

```
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
```



Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

## XGBoost

```
[ ]: from xgboost import XGBClassifier

xgb_cls = XGBClassifier(max_depth = 5, objective = 'reg:logistic',
                        num_parallel_tree = 20, booster = 'gbtree',
                        gamma = 0.5, tree_method = 'gpu_hist', subsample = 0.4, reg_lambda = 1)
xgb_cls.fit(x_train_aug.reshape(n_aug_train, -1), y_train_aug)
preds_train = xgb_cls.predict(x_train_aug.reshape(n_aug_train, -1))
preds_val = xgb_cls.predict(x_val.reshape(n_val, -1))
preds_test = xgb_cls.predict(x_test.reshape(n_test, -1))

print("Training Classification Report: \n", classification_report(y_train_aug, preds_train, digits = N_DIGITS))
print("\nValidation Classification Report: \n", classification_report(y_val, preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test, preds_test, digits = N_DIGITS))
```

### Training Classification Report:

	precision	recall	f1-score	support
0.0	0.915	0.954	0.934	20520
1.0	0.953	0.912	0.932	20682
accuracy			0.933	41202
macro avg	0.934	0.933	0.933	41202
weighted avg	0.934	0.933	0.933	41202

### Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.498	1.000	0.665	1466
1.0	0.000	0.000	0.000	1477
accuracy			0.498	2943
macro avg	0.249	0.500	0.333	2943
weighted avg	0.248	0.498	0.331	2943

### Testing Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

	0.0	0.498	1.000	0.665	1303
	1.0	0.000	0.000	0.000	1314
accuracy				0.498	2617
macro avg	0.249	0.500	0.332		2617
weighted avg	0.248	0.498	0.331		2617

```
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
```

## SVM

```
[ ]: from sklearn.svm import SVC

svm_cls = SVC(kernel = 'rbf', max_iter = 250, verbose= True)
svm_cls.fit(x_train_aug.reshape(n_aug_train, -1), y_train_aug)
preds_train = svm_cls.predict(x_train_aug.reshape(n_aug_train, -1))
preds_val   = svm_cls.predict(x_val.reshape(n_val, -1))
preds_test  = svm_cls.predict(x_test.reshape(n_test, -1))
```

```

print("Training Classification Report: \n", classification_report(y_train_aug,
↳preds_train, digits = N_DIGITS))
print("\nValidation Classification Report: \n", classification_report(y_val,
↳preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test,
↳preds_test, digits = N_DIGITS))

```

[LibSVM]

c:\ProgramData\Anaconda3\envs\ml\lib\site-packages\sklearn\svm\\_base.py:301:  
ConvergenceWarning: Solver terminated early (max\_iter=250). Consider pre-  
processing your data with StandardScaler or MinMaxScaler.  
warnings.warn(

Training Classification Report:

	precision	recall	f1-score	support
0.0	0.449	0.132	0.204	20520
1.0	0.494	0.840	0.622	20682
accuracy			0.487	41202
macro avg	0.472	0.486	0.413	41202
weighted avg	0.472	0.487	0.414	41202

Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.498	1.000	0.665	1466
1.0	0.000	0.000	0.000	1477
accuracy			0.498	2943
macro avg	0.249	0.500	0.333	2943
weighted avg	0.248	0.498	0.331	2943

Testing Classification Report:

	precision	recall	f1-score	support
0.0	0.498	1.000	0.665	1303
1.0	0.000	0.000	0.000	1314
accuracy			0.498	2617
macro avg	0.249	0.500	0.332	2617
weighted avg	0.248	0.498	0.331	2617

```

c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))

```

## Transfer Learning

```

[ ]: imagenet = tf.keras.applications.Xception(
    weights = 'imagenet',
    include_top = False,
    input_shape = (72, 72, 3)
)
imagenet.trainable = False

trans_learn = tf.keras.Sequential([
    tf.keras.layers.Resizing(72, 72),
    imagenet,

    tf.keras.layers.GlobalAveragePooling2D(),
    tf.keras.layers.Dense(512, activation = 'relu'),
    tf.keras.layers.Dense(512, activation = 'relu'),
    tf.keras.layers.Dense(2, activation = 'sigmoid')

```

```

])

trans_learn.compile(
    optimizer = tf.keras.optimizers.Adam(learning_rate=0.001),
    loss = 'sparse_categorical_crossentropy',
    metrics = ['accuracy']
)

history = trans_learn.fit(
    x_train_aug, y_train_aug, batch_size = 128,
    shuffle = True,
    epochs = 50,
    validation_data = [x_val, y_val],
    validation_batch_size = 64,
    callbacks = [tf.keras.callbacks.EarlyStopping(monitor = 'val_loss',
    ↪patience = 5),
                 tf.keras.callbacks.ModelCheckpoint(str = 'val_loss',
    ↪save_best_only = True, save_weights_only = True, filepath= "TL_Aug")]
    )

trans_learn.load_weights('TL_Aug')

fig, ax = plot_history(history)

fig.savefig(PLOT_SAVE + "TL_aug_losscurve.png", facecolor = 'white')
preds_train = np.argmax(trans_learn.predict(x_train_aug), axis = 1)
preds_val    = np.argmax(trans_learn.predict(x_val), axis = 1)
preds_test   = np.argmax(trans_learn.predict(x_test), axis = 1)

print("Training Classification Report: \n", classification_report(y_train_aug,
    ↪preds_train, digits = N_DIGITS))
print("\nValidation Classification Report: \n", classification_report(y_val,
    ↪preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test,
    ↪preds_test, digits = N_DIGITS))

```

Epoch 1/50

322/322 [=====] - 11s 27ms/step - loss: 1.2411 -  
accuracy: 0.6869 - val\_loss: 0.6934 - val\_accuracy: 0.5019

Epoch 2/50

322/322 [=====] - 8s 24ms/step - loss: 0.5090 -  
accuracy: 0.7442 - val\_loss: 0.6945 - val\_accuracy: 0.5019

Epoch 3/50

322/322 [=====] - 8s 24ms/step - loss: 0.4849 -  
accuracy: 0.7584 - val\_loss: 0.6968 - val\_accuracy: 0.5019

Epoch 4/50

322/322 [=====] - 8s 24ms/step - loss: 0.4753 -

accuracy: 0.7656 - val\_loss: 0.6984 - val\_accuracy: 0.5019

Epoch 5/50


322/322 [=====] - 8s 24ms/step - loss: 0.4640 -

accuracy: 0.7715 - val\_loss: 0.7003 - val\_accuracy: 0.5019

Epoch 6/50

322/322 [=====] - 8s 24ms/step - loss: 0.4652 -

accuracy: 0.7737 - val\_loss: 0.7037 - val\_accuracy: 0.5019



1288/1288 [=====] - 9s 7ms/step

92/92 [=====] - 1s 7ms/step

82/82 [=====] - 1s 7ms/step

Training Classification Report:

	precision	recall	f1-score	support
0.0	0.765	0.754	0.760	20520
1.0	0.760	0.770	0.765	20682
accuracy			0.762	41202
macro avg	0.762	0.762	0.762	41202
weighted avg	0.762	0.762	0.762	41202

Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.000	0.000	0.000	1466
1.0	0.502	1.000	0.668	1477

accuracy			0.502	2943
macro avg	0.251	0.500	0.334	2943
weighted avg	0.252	0.502	0.335	2943

#### Testing Classification Report:

	precision	recall	f1-score	support
0.0	0.000	0.000	0.000	1303
1.0	0.502	1.000	0.669	1314

accuracy			0.502	2617
macro avg	0.251	0.500	0.334	2617
weighted avg	0.252	0.502	0.336	2617

```
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
```

## CNN Model

```
[ ]: cnn = tf.keras.Sequential([
    tf.keras.layers.Conv2D(16, (3,3), padding = 'same', activation = 'relu',
    ↪input_shape = x_train[0].shape),
    tf.keras.layers.MaxPool2D((3,3), padding = 'same'),

    tf.keras.layers.Conv2D(32, (2,2), padding = 'same', activation = 'relu'),
    tf.keras.layers.MaxPool2D((2,2), padding = 'same'),

    tf.keras.layers.Flatten(),

    tf.keras.layers.Dense(256, activation = 'relu'),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(256, activation = 'relu'),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(2, activation = 'sigmoid')
])

cnn.compile(
    optimizer = tf.keras.optimizers.Adam(learning_rate=0.001),
    loss = 'sparse_categorical_crossentropy',
    metrics = ['accuracy']
)

history = cnn.fit(
    x_train_aug, y_train_aug, batch_size = 256,
    epochs = 50,
    shuffle = True,
    validation_data = [x_val, y_val],
    validation_batch_size = 128,
    callbacks = [
        tf.keras.callbacks.EarlyStopping(monitor = 'val_loss', patience
    ↪= 5),
        tf.keras.callbacks.ModelCheckpoint(str = 'val_loss',
    ↪save_best_only = True, save_weights_only = True, filepath= "CNN_Aug")
    ]
)

fig, ax = plot_history(history)

cnn.load_weights('CNN_Aug')

fig.savefig(PLOT_SAVE + "CNN_aug_losscurve.png", facecolor = 'white')
preds_train = np.argmax(cnn.predict(x_train_aug), axis = 1)
preds_val    = np.argmax(cnn.predict(x_val), axis = 1)
preds_test   = np.argmax(cnn.predict(x_test), axis = 1)
```



```

print("Training Classification Report: \n", classification_report(y_train_aug,
↳ preds_train, digits = N_DIGITS))
print("\nValidation Classification Report: \n", classification_report(y_val,
↳ preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test,
↳ preds_test, digits = N_DIGITS))

```

Epoch 1/50

161/161 [=====] - 1s 4ms/step - loss: 3.2800 -  
accuracy: 0.5011 - val\_loss: 0.6931 - val\_accuracy: 0.5019

Epoch 2/50

161/161 [=====] - 1s 3ms/step - loss: 0.6945 -  
accuracy: 0.5011 - val\_loss: 0.6932 - val\_accuracy: 0.5019

Epoch 3/50

161/161 [=====] - 1s 3ms/step - loss: 0.6905 -  
accuracy: 0.5182 - val\_loss: 0.6933 - val\_accuracy: 0.4981

Epoch 4/50


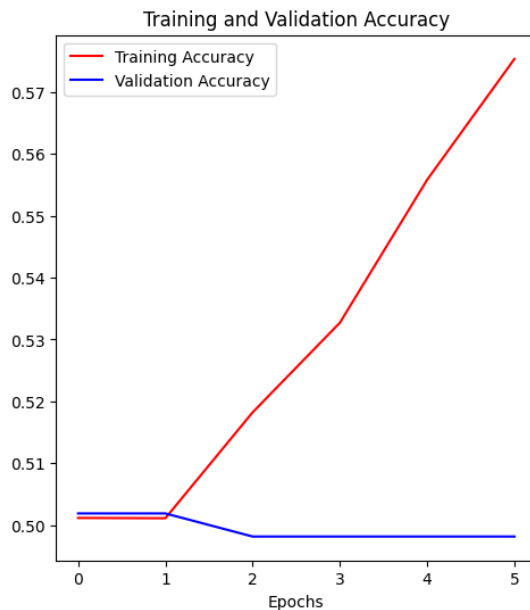
161/161 [=====] - 1s 3ms/step - loss: 0.6869 -  
accuracy: 0.5327 - val\_loss: 0.6934 - val\_accuracy: 0.4981

Epoch 5/50

161/161 [=====] - 1s 3ms/step - loss: 0.6775 -  
accuracy: 0.5558 - val\_loss: 0.6937 - val\_accuracy: 0.4981

Epoch 6/50

161/161 [=====] - 1s 3ms/step - loss: 0.6701 -  
accuracy: 0.5754 - val\_loss: 0.6960 - val\_accuracy: 0.4981



```
1288/1288 [=====] - 1s 1ms/step
92/92 [=====] - 0s 1ms/step
82/82 [=====] - 0s 1ms/step
```

Training Classification Report:

	precision	recall	f1-score	support
0.0	0.487	0.011	0.021	20520
1.0	0.502	0.989	0.666	20682
accuracy			0.502	41202
macro avg	0.494	0.500	0.344	41202
weighted avg	0.494	0.502	0.345	41202

Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.000	0.000	0.000	1466
1.0	0.502	1.000	0.668	1477
accuracy			0.502	2943
macro avg	0.251	0.500	0.334	2943
weighted avg	0.252	0.502	0.335	2943

Testing Classification Report:

	precision	recall	f1-score	support
0.0	0.000	0.000	0.000	1303
1.0	0.502	1.000	0.669	1314
accuracy			0.502	2617
macro avg	0.251	0.500	0.334	2617
weighted avg	0.252	0.502	0.336	2617

```
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
```

Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))
```

## 0.7 HSV Modelling with only the saturation dimension

Transfer learning does not work since it requires exactly 3 input channels (Xception, VGG19, etc)

### 0.7.1 Preparing Datasets

```
[ ]: temp = []
for img in tqdm(x_train):
    temp.append(rgb2hsv(img))
temp = np.array(temp)[..., 1]
x_train = temp
x_train = x_train[..., np.newaxis]

temp = []
for img in tqdm(x_val):
    temp.append(rgb2hsv(img))
temp = np.array(temp)[..., 1]
x_val = temp
x_val = x_val[..., np.newaxis]

temp = []
for img in tqdm(x_test):
    temp.append(rgb2hsv(img))
temp = np.array(temp)[..., 1]
x_test = temp
x_test = x_test[..., np.newaxis]

temp = []
```

```

for img in tqdm(x_train_aug):
    temp.append(rgb2hsv(img))
temp = np.array(temp)[..., 1]
x_train_aug = temp
x_train_aug = x_train_aug[..., np.newaxis]

```

```

x_train.shape, x_val.shape, x_test.shape

```

```

100%|      | 20601/20601 [00:03<00:00, 6839.57it/s]
100%|      | 2943/2943 [00:00<00:00, 6781.57it/s]
100%|      | 2617/2617 [00:00<00:00, 6946.58it/s]
100%|      | 41202/41202 [00:06<00:00, 6203.58it/s]

```

```

[ ]: ((20601, 25, 25, 1), (2943, 25, 25, 1), (2617, 25, 25, 1))

```

## 0.7.2 Unaugmented

### Naive Bayes

```

[ ]: from sklearn import naive_bayes

nb_cls = naive_bayes.GaussianNB()
nb_cls.fit(x_train.reshape(n_train, -1), y_train)
preds_train = nb_cls.predict(x_train.reshape(n_train, -1))
preds_val    = nb_cls.predict(x_val.reshape(n_val, -1))
preds_test   = nb_cls.predict(x_test.reshape(n_test, -1))

print("Training Classification Report: \n", classification_report(y_train,
    ↪preds_train, digits = N_DIGITS))
print("\nValidation Classification Report: \n", classification_report(y_val,
    ↪preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test,
    ↪preds_test, digits = N_DIGITS))

# eval = testing_module.ModelEvaluation(y_val,
#     nb_cls.predict(x_val.reshape(n_val, -1)),
#     model_reference_name = 'Naive Bayes',
#     model_type = 'classification',
#     plot_classification_metric = ['roc_auc']) # if classification
# eval.evaluate(evaluate_save= True, plots_show = True)

```

Training Classification Report:

	precision	recall	f1-score	support
0.0	0.607	0.565	0.585	10260
1.0	0.596	0.636	0.615	10341

accuracy			0.601	20601
macro avg	0.601	0.601	0.600	20601
weighted avg	0.601	0.601	0.600	20601

#### Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.604	0.586	0.595	1466
1.0	0.601	0.618	0.609	1477
accuracy			0.602	2943
macro avg	0.602	0.602	0.602	2943
weighted avg	0.602	0.602	0.602	2943

#### Testing Classification Report:

	precision	recall	f1-score	support
0.0	0.632	0.585	0.607	1303
1.0	0.617	0.662	0.639	1314
accuracy			0.624	2617
macro avg	0.624	0.623	0.623	2617
weighted avg	0.624	0.624	0.623	2617

### Logistic Regression

```
[ ]: from sklearn.linear_model import LogisticRegression

logreg_cls = tf.keras.Sequential([tf.keras.layers.Dense(2, activation =_
    ↳ 'sigmoid', input_dim = HEIGHT * WIDTH)])
logreg_cls.compile(optimizer = tf.optimizers.Adam(0.002), loss =_
    ↳ 'sparse_categorical_crossentropy', metrics = ['accuracy'])
history = logreg_cls.fit(x_train.reshape(n_train, -1),
                        y_train,
                        batch_size = 512,
                        validation_data=[x_val.reshape(n_val, -1), y_val],
                        validation_batch_size=128,
                        epochs = 100,
                        callbacks = [tf.keras.callbacks.EarlyStopping(monitor=_
    ↳ 'val_loss', patience = 10),
                                tf.keras.callbacks.ModelCheckpoint(str =_
    ↳ 'val_loss', save_best_only = True, save_weights_only = True, filepath=_
    ↳ "LR_No_Aug_HSV")],
                        verbose = 1)
```

```

logreg_cls.load_weights('LR_No_Aug_HSV')
preds_train = np.argmax(logreg_cls.predict(x_train.reshape(n_train, -1)), axis=
    ↪ 1)
preds_val    = np.argmax(logreg_cls.predict(x_val.reshape(n_val, -1)), axis = 1)
preds_test   = np.argmax(logreg_cls.predict(x_test.reshape(n_test, -1)), axis =
    ↪ 1)

fig, ax = plot_history(history)

fig.savefig(PLOT_SAVE + "logreg_unaug_losscurve_HSV.png", facecolor = 'white')

print("Training Classification Report: \n", classification_report(y_train,
    ↪ preds_train, digits = N_DIGITS))
print("\nValidation Classification Report: \n", classification_report(y_val,
    ↪ preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test,
    ↪ preds_test, digits = N_DIGITS))

```

Epoch 1/100

41/41 [=====] - 1s 6ms/step - loss: 0.6936 - accuracy: 0.5144 - val\_loss: 0.6743 - val\_accuracy: 0.5260

Epoch 2/100

41/41 [=====] - 0s 3ms/step - loss: 0.6704 - accuracy: 0.5738 - val\_loss: 0.6610 - val\_accuracy: 0.5929

Epoch 3/100

41/41 [=====] - 0s 3ms/step - loss: 0.6579 - accuracy: 0.6304 - val\_loss: 0.6498 - val\_accuracy: 0.6521

Epoch 4/100

41/41 [=====] - 0s 3ms/step - loss: 0.6484 - accuracy: 0.6475 - val\_loss: 0.6431 - val\_accuracy: 0.6677

Epoch 5/100

41/41 [=====] - 0s 3ms/step - loss: 0.6426 - accuracy: 0.6536 - val\_loss: 0.6367 - val\_accuracy: 0.6558

Epoch 6/100

41/41 [=====] - 0s 3ms/step - loss: 0.6360 - accuracy: 0.6585 - val\_loss: 0.6295 - val\_accuracy: 0.6701

Epoch 7/100

41/41 [=====] - 0s 3ms/step - loss: 0.6309 - accuracy: 0.6639 - val\_loss: 0.6249 - val\_accuracy: 0.6765

Epoch 8/100

41/41 [=====] - 0s 3ms/step - loss: 0.6281 - accuracy: 0.6604 - val\_loss: 0.6234 - val\_accuracy: 0.6660

Epoch 9/100

41/41 [=====] - 0s 3ms/step - loss: 0.6248 - accuracy: 0.6643 - val\_loss: 0.6195 - val\_accuracy: 0.6762

Epoch 10/100

41/41 [=====] - 0s 3ms/step - loss: 0.6213 - accuracy:

0.6685 - val\_loss: 0.6161 - val\_accuracy: 0.6704  
 Epoch 11/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6190 - accuracy:  
 0.6676 - val\_loss: 0.6161 - val\_accuracy: 0.6724  
 Epoch 12/100  
 41/41 [=====] - 0s 4ms/step - loss: 0.6192 - accuracy:  
 0.6609 - val\_loss: 0.6135 - val\_accuracy: 0.6721  
 Epoch 13/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6175 - accuracy:  
 0.6661 - val\_loss: 0.6260 - val\_accuracy: 0.6463  
 Epoch 14/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6166 - accuracy:  
 0.6676 - val\_loss: 0.6092 - val\_accuracy: 0.6762  
 Epoch 15/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6146 - accuracy:  
 0.6663 - val\_loss: 0.6089 - val\_accuracy: 0.6680  
 Epoch 16/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6123 - accuracy:  
 0.6716 - val\_loss: 0.6089 - val\_accuracy: 0.6639  
 Epoch 17/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6114 - accuracy:  
 0.6711 - val\_loss: 0.6149 - val\_accuracy: 0.6619  
 Epoch 18/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6115 - accuracy:  
 0.6698 - val\_loss: 0.6057 - val\_accuracy: 0.6762  
 Epoch 19/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6101 - accuracy:  
 0.6714 - val\_loss: 0.6049 - val\_accuracy: 0.6741  
 Epoch 20/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6075 - accuracy:  
 0.6748 - val\_loss: 0.6042 - val\_accuracy: 0.6724  
 Epoch 21/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6077 - accuracy:  
 0.6744 - val\_loss: 0.6047 - val\_accuracy: 0.6789  
 Epoch 22/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6070 - accuracy:  
 0.6755 - val\_loss: 0.6041 - val\_accuracy: 0.6687  
 Epoch 23/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6056 - accuracy:  
 0.6775 - val\_loss: 0.6025 - val\_accuracy: 0.6735  
 Epoch 24/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6058 - accuracy:  
 0.6781 - val\_loss: 0.6038 - val\_accuracy: 0.6769  
 Epoch 25/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6059 - accuracy:  
 0.6766 - val\_loss: 0.6027 - val\_accuracy: 0.6772  
 Epoch 26/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6059 - accuracy:

0.6740 - val\_loss: 0.6086 - val\_accuracy: 0.6728  
 Epoch 27/100  
 41/41 [=====] - 0s 4ms/step - loss: 0.6047 - accuracy:  
 0.6768 - val\_loss: 0.6010 - val\_accuracy: 0.6738  
 Epoch 28/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6033 - accuracy:  
 0.6811 - val\_loss: 0.6004 - val\_accuracy: 0.6731  
 Epoch 29/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6034 - accuracy:  
 0.6794 - val\_loss: 0.6013 - val\_accuracy: 0.6782  
 Epoch 30/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6035 - accuracy:  
 0.6797 - val\_loss: 0.6008 - val\_accuracy: 0.6752  
 Epoch 31/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6026 - accuracy:  
 0.6810 - val\_loss: 0.5995 - val\_accuracy: 0.6799  
 Epoch 32/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6018 - accuracy:  
 0.6822 - val\_loss: 0.5992 - val\_accuracy: 0.6762  
 Epoch 33/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6014 - accuracy:  
 0.6815 - val\_loss: 0.6132 - val\_accuracy: 0.6643  
 Epoch 34/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6028 - accuracy:  
 0.6782 - val\_loss: 0.5985 - val\_accuracy: 0.6792  
 Epoch 35/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6006 - accuracy:  
 0.6839 - val\_loss: 0.5995 - val\_accuracy: 0.6735  
 Epoch 36/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6010 - accuracy:  
 0.6832 - val\_loss: 0.6046 - val\_accuracy: 0.6792  
 Epoch 37/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6009 - accuracy:  
 0.6825 - val\_loss: 0.5978 - val\_accuracy: 0.6806  
 Epoch 38/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5998 - accuracy:  
 0.6845 - val\_loss: 0.5980 - val\_accuracy: 0.6823  
 Epoch 39/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6002 - accuracy:  
 0.6841 - val\_loss: 0.5985 - val\_accuracy: 0.6752  
 Epoch 40/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5993 - accuracy:  
 0.6857 - val\_loss: 0.5987 - val\_accuracy: 0.6714  
 Epoch 41/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5998 - accuracy:  
 0.6830 - val\_loss: 0.5979 - val\_accuracy: 0.6854  
 Epoch 42/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5990 - accuracy:



0.6849 - val\_loss: 0.5970 - val\_accuracy: 0.6796  
 Epoch 43/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.6011 - accuracy:  
 0.6822 - val\_loss: 0.5969 - val\_accuracy: 0.6854  
 Epoch 44/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5980 - accuracy:  
 0.6870 - val\_loss: 0.5965 - val\_accuracy: 0.6823  
 Epoch 45/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5992 - accuracy:  
 0.6856 - val\_loss: 0.6001 - val\_accuracy: 0.6908  
 Epoch 46/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5981 - accuracy:  
 0.6867 - val\_loss: 0.5965 - val\_accuracy: 0.6867  
 Epoch 47/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5997 - accuracy:  
 0.6827 - val\_loss: 0.5963 - val\_accuracy: 0.6820  
 Epoch 48/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5975 - accuracy:  
 0.6870 - val\_loss: 0.5972 - val\_accuracy: 0.6908  
 Epoch 49/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5979 - accuracy:  
 0.6889 - val\_loss: 0.5961 - val\_accuracy: 0.6820  
 Epoch 50/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5987 - accuracy:  
 0.6851 - val\_loss: 0.5975 - val\_accuracy: 0.6741  
 Epoch 51/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5987 - accuracy:  
 0.6856 - val\_loss: 0.5957 - val\_accuracy: 0.6823  
 Epoch 52/100  
 41/41 [=====] - 0s 4ms/step - loss: 0.5981 - accuracy:  
 0.6860 - val\_loss: 0.5952 - val\_accuracy: 0.6830  
 Epoch 53/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5971 - accuracy:  
 0.6860 - val\_loss: 0.5953 - val\_accuracy: 0.6843  
 Epoch 54/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5961 - accuracy:  
 0.6896 - val\_loss: 0.5951 - val\_accuracy: 0.6816  
 Epoch 55/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5970 - accuracy:  
 0.6878 - val\_loss: 0.5958 - val\_accuracy: 0.6867  
 Epoch 56/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5963 - accuracy:  
 0.6903 - val\_loss: 0.6012 - val\_accuracy: 0.6898  
 Epoch 57/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5977 - accuracy:  
 0.6855 - val\_loss: 0.5999 - val\_accuracy: 0.6932  
 Epoch 58/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5964 - accuracy:


0.6895 - val\_loss: 0.5950 - val\_accuracy: 0.6850  
 Epoch 59/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5971 - accuracy:  
 0.6866 - val\_loss: 0.5945 - val\_accuracy: 0.6826  
 Epoch 60/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5959 - accuracy:  
 0.6888 - val\_loss: 0.5990 - val\_accuracy: 0.6942  
 Epoch 61/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5965 - accuracy:  
 0.6897 - val\_loss: 0.5986 - val\_accuracy: 0.6745  
 Epoch 62/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5970 - accuracy:  
 0.6876 - val\_loss: 0.5959 - val\_accuracy: 0.6925  
 Epoch 63/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5954 - accuracy:  
 0.6916 - val\_loss: 0.5945 - val\_accuracy: 0.6857  
 Epoch 64/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5972 - accuracy:  
 0.6876 - val\_loss: 0.5964 - val\_accuracy: 0.6775  
 Epoch 65/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5953 - accuracy:  
 0.6902 - val\_loss: 0.5944 - val\_accuracy: 0.6833  
 Epoch 66/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5949 - accuracy:  
 0.6917 - val\_loss: 0.5953 - val\_accuracy: 0.6789  
 Epoch 67/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5950 - accuracy:  
 0.6924 - val\_loss: 0.5950 - val\_accuracy: 0.6789  
 Epoch 68/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5949 - accuracy:  
 0.6890 - val\_loss: 0.5954 - val\_accuracy: 0.6918  
 Epoch 69/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5951 - accuracy:  
 0.6932 - val\_loss: 0.5954 - val\_accuracy: 0.6935  
 Epoch 70/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5946 - accuracy:  
 0.6933 - val\_loss: 0.5950 - val\_accuracy: 0.6786  
 Epoch 71/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5939 - accuracy:  
 0.6925 - val\_loss: 0.5953 - val\_accuracy: 0.6938  
 Epoch 72/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5938 - accuracy:  
 0.6946 - val\_loss: 0.5938 - val\_accuracy: 0.6843  
 Epoch 73/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5945 - accuracy:  
 0.6922 - val\_loss: 0.5935 - val\_accuracy: 0.6830  
 Epoch 74/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5932 - accuracy:

0.6955 - val\_loss: 0.5949 - val\_accuracy: 0.6925  
 Epoch 75/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5937 - accuracy:  
 0.6926 - val\_loss: 0.5937 - val\_accuracy: 0.6840  
 Epoch 76/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5942 - accuracy:  
 0.6929 - val\_loss: 0.5982 - val\_accuracy: 0.6955  
 Epoch 77/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5947 - accuracy:  
 0.6891 - val\_loss: 0.5950 - val\_accuracy: 0.6945  
 Epoch 78/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5945 - accuracy:  
 0.6889 - val\_loss: 0.5937 - val\_accuracy: 0.6874  
 Epoch 79/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5938 - accuracy:  
 0.6911 - val\_loss: 0.5931 - val\_accuracy: 0.6837  
 Epoch 80/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5935 - accuracy:  
 0.6935 - val\_loss: 0.5930 - val\_accuracy: 0.6833  
 Epoch 81/100  
 41/41 [=====] - 0s 4ms/step - loss: 0.5931 - accuracy:  
 0.6934 - val\_loss: 0.5942 - val\_accuracy: 0.6908  
 Epoch 82/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5928 - accuracy:  
 0.6949 - val\_loss: 0.5931 - val\_accuracy: 0.6840  
 Epoch 83/100  
 41/41 [=====] - 0s 4ms/step - loss: 0.5940 - accuracy:  
 0.6904 - val\_loss: 0.5945 - val\_accuracy: 0.6922  
 Epoch 84/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5928 - accuracy:  
 0.6964 - val\_loss: 0.5930 - val\_accuracy: 0.6884  
 Epoch 85/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5957 - accuracy:  
 0.6882 - val\_loss: 0.5945 - val\_accuracy: 0.6949  
 Epoch 86/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5937 - accuracy:  
 0.6925 - val\_loss: 0.5931 - val\_accuracy: 0.6857  
 Epoch 87/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5924 - accuracy:  
 0.6940 - val\_loss: 0.5932 - val\_accuracy: 0.6888  
 Epoch 88/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5931 - accuracy:  
 0.6935 - val\_loss: 0.5929 - val\_accuracy: 0.6857  
 Epoch 89/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5936 - accuracy:  
 0.6940 - val\_loss: 0.5935 - val\_accuracy: 0.6884  
 Epoch 90/100  
 41/41 [=====] - 0s 3ms/step - loss: 0.5924 - accuracy:

```

0.6942 - val_loss: 0.5972 - val_accuracy: 0.6979
Epoch 91/100
41/41 [=====] - 0s 3ms/step - loss: 0.5935 - accuracy:
0.6943 - val_loss: 0.5958 - val_accuracy: 0.6976
Epoch 92/100
41/41 [=====] - 0s 3ms/step - loss: 0.5925 - accuracy:
0.6935 - val_loss: 0.5929 - val_accuracy: 0.6847
Epoch 93/100
41/41 [=====] - 0s 3ms/step - loss: 0.5935 - accuracy:
0.6940 - val_loss: 0.5951 - val_accuracy: 0.6986
Epoch 94/100
41/41 [=====] - 0s 3ms/step - loss: 0.5920 - accuracy:
0.6949 - val_loss: 0.5962 - val_accuracy: 0.6993
Epoch 95/100
41/41 [=====] - 0s 3ms/step - loss: 0.5939 - accuracy:
0.6920 - val_loss: 0.5940 - val_accuracy: 0.6952
Epoch 96/100
41/41 [=====] - 0s 3ms/step - loss: 0.5918 - accuracy:
0.6939 - val_loss: 0.5946 - val_accuracy: 0.6955
Epoch 97/100
41/41 [=====] - 0s 3ms/step - loss: 0.5924 - accuracy:
0.6925 - val_loss: 0.5982 - val_accuracy: 0.6932
Epoch 98/100
41/41 [=====] - 0s 3ms/step - loss: 0.5926 - accuracy:
0.6937 - val_loss: 0.5924 - val_accuracy: 0.6857
Epoch 99/100
41/41 [=====] - 0s 3ms/step - loss: 0.5911 - accuracy:
0.6960 - val_loss: 0.5984 - val_accuracy: 0.6952
Epoch 100/100
41/41 [=====] - 0s 4ms/step - loss: 0.5920 - accuracy:
0.6960 - val_loss: 0.5922 - val_accuracy: 0.6867
644/644 [=====] - 1s 719us/step
92/92 [=====] - 0s 709us/step
82/82 [=====] - 0s 871us/step

```



#### Training Classification Report:

	precision	recall	f1-score	support
0.0	0.692	0.706	0.699	10260
1.0	0.702	0.688	0.695	10341
accuracy			0.697	20601
macro avg	0.697	0.697	0.697	20601
weighted avg	0.697	0.697	0.697	20601

#### Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.679	0.703	0.691	1466
1.0	0.694	0.671	0.683	1477
accuracy			0.687	2943
macro avg	0.687	0.687	0.687	2943
weighted avg	0.687	0.687	0.687	2943

#### Testing Classification Report:

	precision	recall	f1-score	support
0.0	0.671	0.687	0.679	1303
1.0	0.682	0.667	0.674	1314

accuracy			0.677	2617
macro avg	0.677	0.677	0.677	2617
weighted avg	0.677	0.677	0.677	2617

## Decision Trees

```
[ ]: from sklearn.tree import DecisionTreeClassifier

dt_cls = DecisionTreeClassifier(criterion = 'gini', splitter = 'best',
    ↪max_depth = 12, min_samples_split = 2, )
dt_cls.fit(x_train.reshape(n_train, -1), y_train)
preds_train = dt_cls.predict(x_train.reshape(n_train, -1))
preds_val    = dt_cls.predict(x_val.reshape(n_val, -1))
preds_test   = dt_cls.predict(x_test.reshape(n_test, -1))

print("Training Classification Report: \n", classification_report(y_train,
    ↪preds_train, digits = N_DIGITS))
print("\nValidation Classification Report: \n", classification_report(y_val,
    ↪preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test,
    ↪preds_test, digits = N_DIGITS))

# eval = testing_module.ModelEvaluation(y_val,
#     nb_cls.predict(x_val.reshape(n_val, -1)),
#     model_reference_name = 'Naive Bayes',
#     model_type = 'classification',
#     plot_classification_metric = ['roc_auc']) # if classification
# eval.evaluate(evaluate_save= True, plots_show = True)
```

### Training Classification Report:

	precision	recall	f1-score	support
0.0	0.761	0.837	0.797	10260
1.0	0.821	0.739	0.778	10341
accuracy			0.788	20601
macro avg	0.791	0.788	0.787	20601
weighted avg	0.791	0.788	0.787	20601

### Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.667	0.760	0.711	1466
1.0	0.724	0.624	0.670	1477
accuracy			0.692	2943

macro avg	0.696	0.692	0.690	2943
weighted avg	0.696	0.692	0.690	2943

Testing Classification Report:

	precision	recall	f1-score	support
0.0	0.653	0.740	0.694	1303
1.0	0.703	0.610	0.653	1314
accuracy			0.675	2617
macro avg	0.678	0.675	0.674	2617
weighted avg	0.678	0.675	0.673	2617

## XGBoost

```
[ ]: tf.keras.backend.clear_session()
from xgboost import XGBClassifier

xgb_cls = XGBClassifier(max_depth = 10, objective = 'reg:logistic',
                        num_parallel_tree = 20, booster = 'gbtree',
                        gamma = 0.5, tree_method = 'gpu_hist', subsample = 0.4, reg_lambda = 1)
setattr(xgb_cls, 'verbosity', 1)
xgb_cls.fit(x_train.reshape(n_train, -1), y_train)
preds_train = xgb_cls.predict(x_train.reshape(n_train, -1))
preds_val = xgb_cls.predict(x_val.reshape(n_val, -1))
preds_test = xgb_cls.predict(x_test.reshape(n_test, -1))

print("Training Classification Report: \n", classification_report(y_train,
    preds_train, digits = N_DIGITS))
print("\nValidation Classification Report: \n", classification_report(y_val,
    preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test,
    preds_test, digits = N_DIGITS))
```

Training Classification Report:

	precision	recall	f1-score	support
0.0	0.998	0.999	0.999	10260
1.0	0.999	0.998	0.999	10341
accuracy			0.999	20601
macro avg	0.999	0.999	0.999	20601
weighted avg	0.999	0.999	0.999	20601

#### Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.904	0.954	0.928	1466
1.0	0.951	0.899	0.924	1477
accuracy			0.926	2943
macro avg	0.927	0.926	0.926	2943
weighted avg	0.928	0.926	0.926	2943

#### Testing Classification Report:

	precision	recall	f1-score	support
0.0	0.920	0.959	0.939	1303
1.0	0.957	0.918	0.937	1314
accuracy			0.938	2617
macro avg	0.939	0.938	0.938	2617
weighted avg	0.939	0.938	0.938	2617

## SVM

```
[ ]: from sklearn.svm import SVC

svm_cls = SVC(kernel = 'poly', degree = 3, gamma = 'auto', max_iter = 250,
    verbose= True)
svm_cls.fit(x_train.reshape(n_train, -1), y_train)
preds_train = svm_cls.predict(x_train.reshape(n_train, -1))
preds_val   = svm_cls.predict(x_val.reshape(n_val, -1))
preds_test  = svm_cls.predict(x_test.reshape(n_test, -1))

print("Training Classification Report: \n", classification_report(y_train,
    preds_train, digits = N_DIGITS))
print("\nValidation Classification Report: \n", classification_report(y_val,
    preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test,
    preds_test, digits = N_DIGITS))
```

[LibSVM]

```
c:\ProgramData\Anaconda3\envs\ml\lib\site-packages\sklearn\svm\_base.py:301:
ConvergenceWarning: Solver terminated early (max_iter=250). Consider pre-
processing your data with StandardScaler or MinMaxScaler.
warnings.warn(
```

#### Training Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------



	0.0	0.200	0.000	0.000	10260
	1.0	0.502	1.000	0.668	10341
accuracy				0.502	20601
macro avg	0.351	0.500	0.334		20601
weighted avg	0.352	0.502	0.336		20601

#### Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.000	0.000	0.000	1466
1.0	0.502	1.000	0.668	1477
accuracy			0.502	2943
macro avg	0.251	0.500	0.334	2943
weighted avg	0.252	0.502	0.335	2943

#### Testing Classification Report:

	precision	recall	f1-score	support
0.0	0.000	0.000	0.000	1303
1.0	0.502	0.999	0.668	1314
accuracy			0.502	2617
macro avg	0.251	0.500	0.334	2617
weighted avg	0.252	0.502	0.336	2617

```
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
c:\ProgramData\Anaconda3\envs\ml\lib\site-
packages\sklearn\metrics\_classification.py:1334: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

#### CNN Model

```

[ ]: cnn = tf.keras.Sequential([
    tf.keras.layers.Conv2D(64, (3,3), padding = 'same', activation = 'relu',
    ↪input_shape = x_train[0].shape),
    tf.keras.layers.MaxPool2D((3,3), padding = 'same'),

    tf.keras.layers.UpSampling2D((4,4)),

    tf.keras.layers.Conv2D(32, (3,3), padding = 'same', activation = 'relu'),
    tf.keras.layers.MaxPool2D((3,3), padding = 'same'),

    tf.keras.layers.Conv2D(32, (2,2), padding = 'same', activation = 'relu'),
    tf.keras.layers.MaxPool2D((2,2), padding = 'same'),

    tf.keras.layers.Flatten(),

    tf.keras.layers.Dense(512, activation = 'relu'),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(512, activation = 'relu'),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(2, activation = 'sigmoid')
])

cnn.compile(
    optimizer = tf.keras.optimizers.Adam(learning_rate=0.003),
    loss = 'sparse_categorical_crossentropy',
    metrics = ['accuracy']
)

history = cnn.fit(
    x_train, y_train, batch_size = 256,
    epochs = 50,
    validation_data = [x_val, y_val],
    validation_batch_size = 256,
    callbacks = [tf.keras.callbacks.EarlyStopping(monitor = 'val_loss',
    ↪patience = 10),
                 tf.keras.callbacks.ModelCheckpoint(str = 'val_loss',
    ↪save_best_only = True, save_weights_only = True, filepath= "CNN_No_Aug_HSV")]
    )

cnn.load_weights('CNN_No_Aug_HSV')
fig, ax = plot_history(history)

fig.savefig(PLOT_SAVE + "CNN_unaug_losscurve_HSV.png", facecolor = 'white')
preds_train = np.argmax(cnn.predict(x_train), axis = 1)
preds_val    = np.argmax(cnn.predict(x_val), axis = 1)

```

```

preds_test = np.argmax(cnn.predict(x_test), axis = 1)

print("Training Classification Report: \n", classification_report(y_train,
↳ preds_train, digits = N_DIGITS))
print("\nValidation Classification Report: \n", classification_report(y_val,
↳ preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test,
↳ preds_test, digits = N_DIGITS))

```

Epoch 1/50

81/81 [=====] - 3s 17ms/step - loss: 0.2080 - accuracy:  
0.9079 - val\_loss: 0.0381 - val\_accuracy: 0.9895

Epoch 2/50

81/81 [=====] - 1s 12ms/step - loss: 0.0343 - accuracy:  
0.9904 - val\_loss: 0.0276 - val\_accuracy: 0.9918

Epoch 3/50

81/81 [=====] - 1s 12ms/step - loss: 0.0274 - accuracy:  
0.9924 - val\_loss: 0.0347 - val\_accuracy: 0.9884

Epoch 4/50

81/81 [=====] - 1s 11ms/step - loss: 0.0329 - accuracy:  
0.9916 - val\_loss: 0.0309 - val\_accuracy: 0.9884

Epoch 5/50

81/81 [=====] - 1s 12ms/step - loss: 0.0237 - accuracy:  
0.9934 - val\_loss: 0.0288 - val\_accuracy: 0.9905

Epoch 6/50

81/81 [=====] - 1s 12ms/step - loss: 0.0221 - accuracy:  
0.9935 - val\_loss: 0.0278 - val\_accuracy: 0.9912

Epoch 7/50

81/81 [=====] - 1s 11ms/step - loss: 0.0205 - accuracy:  
0.9935 - val\_loss: 0.0346 - val\_accuracy: 0.9918

Epoch 8/50

81/81 [=====] - 1s 11ms/step - loss: 0.0210 - accuracy:  
0.9942 - val\_loss: 0.0285 - val\_accuracy: 0.9925

Epoch 9/50

81/81 [=====] - 1s 12ms/step - loss: 0.0207 - accuracy:  
0.9933 - val\_loss: 0.0288 - val\_accuracy: 0.9901

Epoch 10/50

81/81 [=====] - 1s 11ms/step - loss: 0.0214 - accuracy:  
0.9941 - val\_loss: 0.0318 - val\_accuracy: 0.9905

Epoch 11/50

81/81 [=====] - 1s 11ms/step - loss: 0.0204 - accuracy:  
0.9939 - val\_loss: 0.0319 - val\_accuracy: 0.9884


Epoch 12/50

81/81 [=====] - 1s 12ms/step - loss: 0.0197 - accuracy:  
0.9939 - val\_loss: 0.0275 - val\_accuracy: 0.9918

Epoch 13/50

81/81 [=====] - 1s 11ms/step - loss: 0.0208 - accuracy:

0.9935 - val\_loss: 0.0352 - val\_accuracy: 0.9901  
Epoch 14/50  
81/81 [=====] - 1s 11ms/step - loss: 0.0162 - accuracy:  
0.9954 - val\_loss: 0.0374 - val\_accuracy: 0.9908  
Epoch 15/50  
81/81 [=====] - 1s 12ms/step - loss: 0.0187 - accuracy:  
0.9944 - val\_loss: 0.0328 - val\_accuracy: 0.9905  
Epoch 16/50  
81/81 [=====] - 1s 11ms/step - loss: 0.0166 - accuracy:  
0.9950 - val\_loss: 0.0364 - val\_accuracy: 0.9912  
Epoch 17/50  
81/81 [=====] - 1s 12ms/step - loss: 0.0159 - accuracy:  
0.9952 - val\_loss: 0.0278 - val\_accuracy: 0.9915  
Epoch 18/50  
81/81 [=====] - 1s 11ms/step - loss: 0.0152 - accuracy:  
0.9957 - val\_loss: 0.0314 - val\_accuracy: 0.9918  
Epoch 19/50  
81/81 [=====] - 1s 11ms/step - loss: 0.0148 - accuracy:  
0.9960 - val\_loss: 0.0304 - val\_accuracy: 0.9912  
Epoch 20/50  
81/81 [=====] - 1s 11ms/step - loss: 0.0130 - accuracy:  
0.9963 - val\_loss: 0.0331 - val\_accuracy: 0.9915  
Epoch 21/50  
81/81 [=====] - 1s 12ms/step - loss: 0.0120 - accuracy:  
0.9960 - val\_loss: 0.0330 - val\_accuracy: 0.9922  
Epoch 22/50  
81/81 [=====] - 1s 11ms/step - loss: 0.0128 - accuracy:  
0.9954 - val\_loss: 0.0309 - val\_accuracy: 0.9929



644/644 [=====] - 1s 1ms/step

92/92 [=====] - 0s 2ms/step

82/82 [=====] - 0s 1ms/step

Training Classification Report:

	precision	recall	f1-score	support
0.0	0.996	0.995	0.996	10260
1.0	0.995	0.996	0.996	10341
accuracy			0.996	20601
macro avg	0.996	0.996	0.996	20601
weighted avg	0.996	0.996	0.996	20601

Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.993	0.990	0.992	1466
1.0	0.991	0.993	0.992	1477
accuracy			0.992	2943
macro avg	0.992	0.992	0.992	2943
weighted avg	0.992	0.992	0.992	2943

Testing Classification Report:

	precision	recall	f1-score	support
0.0	0.994	0.992	0.993	1303
1.0	0.992	0.994	0.993	1314
accuracy			0.993	2617
macro avg	0.993	0.993	0.993	2617
weighted avg	0.993	0.993	0.993	2617

### 0.7.3 Augmented Dataset

#### Naive Bayes

```
[ ]: from sklearn import naive_bayes

nb_cls = naive_bayes.GaussianNB()
nb_cls.fit(x_train_aug.reshape(n_aug_train, -1), y_train_aug)
preds_train = nb_cls.predict(x_train_aug.reshape(n_aug_train, -1))
preds_val   = nb_cls.predict(x_val.reshape(n_val, -1))
preds_test  = nb_cls.predict(x_test.reshape(n_test, -1))
```

```
print("Training Classification Report: \n", classification_report(y_train_aug,
↳ preds_train, digits = N_DIGITS))
print("\nValidation Classification Report: \n", classification_report(y_val,
↳ preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test,
↳ preds_test, digits = N_DIGITS))
```

#### Training Classification Report:

	precision	recall	f1-score	support
0.0	0.607	0.612	0.609	20520
1.0	0.612	0.607	0.609	20682
accuracy			0.609	41202
macro avg	0.609	0.609	0.609	41202
weighted avg	0.609	0.609	0.609	41202

#### Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.610	0.493	0.545	1466
1.0	0.577	0.687	0.627	1477
accuracy			0.590	2943
macro avg	0.593	0.590	0.586	2943
weighted avg	0.593	0.590	0.586	2943

#### Testing Classification Report:

	precision	recall	f1-score	support
0.0	0.653	0.495	0.563	1303
1.0	0.596	0.739	0.660	1314
accuracy			0.618	2617
macro avg	0.624	0.617	0.611	2617
weighted avg	0.624	0.618	0.612	2617

### Logistic Regression

```
[ ]: from sklearn.linear_model import LogisticRegression

logreg_cls = tf.keras.Sequential([tf.keras.layers.Dense(2, activation =
↳ 'sigmoid', input_dim = HEIGHT * WIDTH)])
```

```

logreg_cls.compile(optimizer = tf.optimizers.Adam(0.0001), loss =_
    ↪ 'sparse_categorical_crossentropy', metrics = ['accuracy'])
history = logreg_cls.fit(x_train_aug.reshape(n_aug_train, -1),
                        y_train_aug,
                        batch_size = 512,
                        validation_data=[x_val.reshape(n_val, -1), y_val],
                        validation_batch_size=128,
                        epochs = 100,
                        callbacks = [tf.keras.callbacks.EarlyStopping(monitor=_
    ↪ 'val_loss', patience = 10),
                                tf.keras.callbacks.ModelCheckpoint(str =_
    ↪ 'val_loss', save_best_only = True, save_weights_only = True, filepath=_
    ↪ "LR_Aug_HSV")])
preds_train = np.argmax(logreg_cls.predict(x_train_aug.reshape(n_aug_train,_
    ↪ -1)), axis = 1)
preds_val    = np.argmax(logreg_cls.predict(x_val.reshape(n_val, -1)), axis = 1)
preds_test   = np.argmax(logreg_cls.predict(x_test.reshape(n_test, -1)), axis =_
    ↪ 1)

logreg_cls.load_weights('LR_Aug_HSV')
fig, ax = plot_history(history)

fig.savefig(PLOT_SAVE + "logreg_aug_losscurve_HSV.png", facecolor = 'white')

print("Training Classification Report: \n", classification_report(y_train_aug,_
    ↪ preds_train, digits = N_DIGITS))
print("\nValidation Classification Report: \n", classification_report(y_val,_
    ↪ preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test,_
    ↪ preds_test, digits = N_DIGITS))

```

Epoch 1/100

81/81 [=====] - 0s 3ms/step - loss: 0.6949 - accuracy:  
0.4968 - val\_loss: 0.6917 - val\_accuracy: 0.4910

Epoch 2/100

81/81 [=====] - 0s 2ms/step - loss: 0.6887 - accuracy:  
0.4986 - val\_loss: 0.6852 - val\_accuracy: 0.4937

Epoch 3/100

81/81 [=====] - 0s 2ms/step - loss: 0.6855 - accuracy:  
0.5055 - val\_loss: 0.6825 - val\_accuracy: 0.4975

Epoch 4/100

81/81 [=====] - 0s 2ms/step - loss: 0.6822 - accuracy:  
0.5168 - val\_loss: 0.6801 - val\_accuracy: 0.5158

Epoch 5/100

81/81 [=====] - 0s 2ms/step - loss: 0.6792 - accuracy:  
0.5251 - val\_loss: 0.6773 - val\_accuracy: 0.5304

Epoch 6/100

81/81 [=====] - 0s 2ms/step - loss: 0.6763 - accuracy: 0.5403 - val\_loss: 0.6743 - val\_accuracy: 0.5375  
Epoch 7/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6736 - accuracy: 0.5521 - val\_loss: 0.6727 - val\_accuracy: 0.5627  
Epoch 8/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6713 - accuracy: 0.5638 - val\_loss: 0.6701 - val\_accuracy: 0.5657  
Epoch 9/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6690 - accuracy: 0.5757 - val\_loss: 0.6682 - val\_accuracy: 0.5691  
Epoch 10/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6668 - accuracy: 0.5844 - val\_loss: 0.6665 - val\_accuracy: 0.5797  
Epoch 11/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6649 - accuracy: 0.5925 - val\_loss: 0.6648 - val\_accuracy: 0.5841  
Epoch 12/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6631 - accuracy: 0.6009 - val\_loss: 0.6634 - val\_accuracy: 0.6018  
Epoch 13/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6614 - accuracy: 0.6089 - val\_loss: 0.6623 - val\_accuracy: 0.5841  
Epoch 14/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6597 - accuracy: 0.6125 - val\_loss: 0.6605 - val\_accuracy: 0.6082  
Epoch 15/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6582 - accuracy: 0.6184 - val\_loss: 0.6594 - val\_accuracy: 0.6086  
Epoch 16/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6568 - accuracy: 0.6232 - val\_loss: 0.6580 - val\_accuracy: 0.6140  
Epoch 17/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6555 - accuracy: 0.6254 - val\_loss: 0.6568 - val\_accuracy: 0.6140  
Epoch 18/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6543 - accuracy: 0.6281 - val\_loss: 0.6560 - val\_accuracy: 0.6140  
Epoch 19/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6531 - accuracy: 0.6299 - val\_loss: 0.6547 - val\_accuracy: 0.6171  
Epoch 20/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6519 - accuracy: 0.6340 - val\_loss: 0.6548 - val\_accuracy: 0.6130  
Epoch 21/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6509 - accuracy: 0.6341 - val\_loss: 0.6533 - val\_accuracy: 0.6154  
Epoch 22/100



81/81 [=====] - 0s 2ms/step - loss: 0.6498 - accuracy: 0.6377 - val\_loss: 0.6530 - val\_accuracy: 0.6154  
Epoch 23/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6489 - accuracy: 0.6368 - val\_loss: 0.6512 - val\_accuracy: 0.6222  
Epoch 24/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6479 - accuracy: 0.6392 - val\_loss: 0.6500 - val\_accuracy: 0.6262  
Epoch 25/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6471 - accuracy: 0.6410 - val\_loss: 0.6504 - val\_accuracy: 0.6194  
Epoch 26/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6463 - accuracy: 0.6410 - val\_loss: 0.6485 - val\_accuracy: 0.6269  
Epoch 27/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6454 - accuracy: 0.6412 - val\_loss: 0.6498 - val\_accuracy: 0.6188  
Epoch 28/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6447 - accuracy: 0.6429 - val\_loss: 0.6478 - val\_accuracy: 0.6239  
Epoch 29/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6439 - accuracy: 0.6436 - val\_loss: 0.6466 - val\_accuracy: 0.6256  
Epoch 30/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6432 - accuracy: 0.6412 - val\_loss: 0.6460 - val\_accuracy: 0.6259  
Epoch 31/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6425 - accuracy: 0.6449 - val\_loss: 0.6458 - val\_accuracy: 0.6276  
Epoch 32/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6418 - accuracy: 0.6427 - val\_loss: 0.6444 - val\_accuracy: 0.6266  
Epoch 33/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6411 - accuracy: 0.6454 - val\_loss: 0.6454 - val\_accuracy: 0.6252  
Epoch 34/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6407 - accuracy: 0.6433 - val\_loss: 0.6440 - val\_accuracy: 0.6252  
Epoch 35/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6399 - accuracy: 0.6448 - val\_loss: 0.6432 - val\_accuracy: 0.6256  
Epoch 36/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6394 - accuracy: 0.6455 - val\_loss: 0.6426 - val\_accuracy: 0.6262  
Epoch 37/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6388 - accuracy: 0.6459 - val\_loss: 0.6432 - val\_accuracy: 0.6252  
Epoch 38/100

81/81 [=====] - 0s 3ms/step - loss: 0.6382 - accuracy:  
0.6466 - val\_loss: 0.6424 - val\_accuracy: 0.6245  
Epoch 39/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6377 - accuracy:  
0.6456 - val\_loss: 0.6411 - val\_accuracy: 0.6269  
Epoch 40/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6373 - accuracy:  
0.6459 - val\_loss: 0.6406 - val\_accuracy: 0.6279  
Epoch 41/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6368 - accuracy:  
0.6437 - val\_loss: 0.6399 - val\_accuracy: 0.6340  
Epoch 42/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6364 - accuracy:  
0.6465 - val\_loss: 0.6400 - val\_accuracy: 0.6262  
Epoch 43/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6358 - accuracy:  
0.6459 - val\_loss: 0.6397 - val\_accuracy: 0.6262  
Epoch 44/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6354 - accuracy:  
0.6455 - val\_loss: 0.6396 - val\_accuracy: 0.6279  
Epoch 45/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6350 - accuracy:  
0.6463 - val\_loss: 0.6394 - val\_accuracy: 0.6276  
Epoch 46/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6345 - accuracy:  
0.6467 - val\_loss: 0.6387 - val\_accuracy: 0.6283  
Epoch 47/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6341 - accuracy:  
0.6465 - val\_loss: 0.6383 - val\_accuracy: 0.6276  
Epoch 48/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6338 - accuracy:  
0.6459 - val\_loss: 0.6376 - val\_accuracy: 0.6290  
Epoch 49/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6333 - accuracy:  
0.6479 - val\_loss: 0.6397 - val\_accuracy: 0.6228  
Epoch 50/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6334 - accuracy:  
0.6474 - val\_loss: 0.6383 - val\_accuracy: 0.6256  
Epoch 51/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6327 - accuracy:  
0.6474 - val\_loss: 0.6381 - val\_accuracy: 0.6245  
Epoch 52/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6323 - accuracy:  
0.6468 - val\_loss: 0.6363 - val\_accuracy: 0.6293  
Epoch 53/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6318 - accuracy:  
0.6484 - val\_loss: 0.6359 - val\_accuracy: 0.6279  
Epoch 54/100


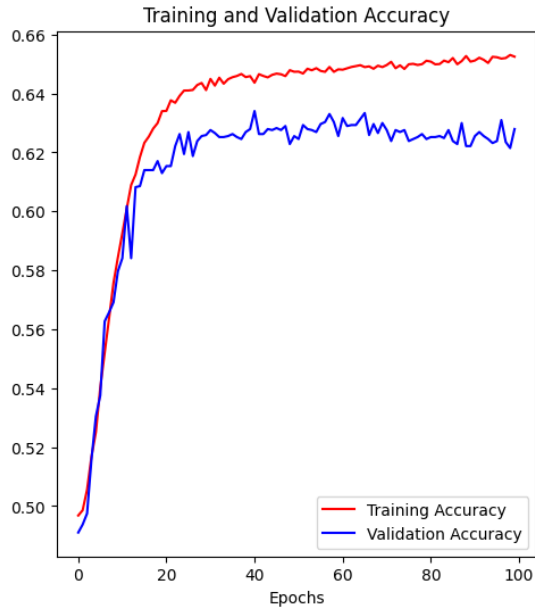
81/81 [=====] - 0s 3ms/step - loss: 0.6316 - accuracy: 0.6479 - val\_loss: 0.6359 - val\_accuracy: 0.6276  
Epoch 55/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6313 - accuracy: 0.6486 - val\_loss: 0.6359 - val\_accuracy: 0.6269  
Epoch 56/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6309 - accuracy: 0.6477 - val\_loss: 0.6352 - val\_accuracy: 0.6296  
Epoch 57/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6306 - accuracy: 0.6473 - val\_loss: 0.6348 - val\_accuracy: 0.6303  
Epoch 58/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6303 - accuracy: 0.6490 - val\_loss: 0.6344 - val\_accuracy: 0.6330  
Epoch 59/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6300 - accuracy: 0.6473 - val\_loss: 0.6342 - val\_accuracy: 0.6303  
Epoch 60/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6296 - accuracy: 0.6482 - val\_loss: 0.6347 - val\_accuracy: 0.6256  
Epoch 61/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6294 - accuracy: 0.6481 - val\_loss: 0.6338 - val\_accuracy: 0.6317  
Epoch 62/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6291 - accuracy: 0.6485 - val\_loss: 0.6336 - val\_accuracy: 0.6290  
Epoch 63/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6288 - accuracy: 0.6490 - val\_loss: 0.6336 - val\_accuracy: 0.6293  
Epoch 64/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6286 - accuracy: 0.6493 - val\_loss: 0.6331 - val\_accuracy: 0.6293  
Epoch 65/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6283 - accuracy: 0.6496 - val\_loss: 0.6329 - val\_accuracy: 0.6313  
Epoch 66/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6281 - accuracy: 0.6489 - val\_loss: 0.6327 - val\_accuracy: 0.6334  
Epoch 67/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6279 - accuracy: 0.6492 - val\_loss: 0.6334 - val\_accuracy: 0.6259  
Epoch 68/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6276 - accuracy: 0.6484 - val\_loss: 0.6323 - val\_accuracy: 0.6296  
Epoch 69/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6274 - accuracy: 0.6494 - val\_loss: 0.6331 - val\_accuracy: 0.6266  
Epoch 70/100

81/81 [=====] - 0s 2ms/step - loss: 0.6271 - accuracy:  
0.6489 - val\_loss: 0.6323 - val\_accuracy: 0.6300  
Epoch 71/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6269 - accuracy:  
0.6497 - val\_loss: 0.6322 - val\_accuracy: 0.6276  
Epoch 72/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6267 - accuracy:  
0.6507 - val\_loss: 0.6330 - val\_accuracy: 0.6239  
Epoch 73/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6265 - accuracy:  
0.6486 - val\_loss: 0.6322 - val\_accuracy: 0.6276  
Epoch 74/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6264 - accuracy:  
0.6496 - val\_loss: 0.6321 - val\_accuracy: 0.6269  
Epoch 75/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6263 - accuracy:  
0.6483 - val\_loss: 0.6319 - val\_accuracy: 0.6276  
Epoch 76/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6260 - accuracy:  
0.6499 - val\_loss: 0.6326 - val\_accuracy: 0.6239  
Epoch 77/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6259 - accuracy:  
0.6501 - val\_loss: 0.6321 - val\_accuracy: 0.6245  
Epoch 78/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6257 - accuracy:  
0.6497 - val\_loss: 0.6314 - val\_accuracy: 0.6252  
Epoch 79/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6254 - accuracy:  
0.6499 - val\_loss: 0.6310 - val\_accuracy: 0.6262  
Epoch 80/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6253 - accuracy:  
0.6511 - val\_loss: 0.6307 - val\_accuracy: 0.6245  
Epoch 81/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6251 - accuracy:  
0.6508 - val\_loss: 0.6313 - val\_accuracy: 0.6252  
Epoch 82/100  
81/81 [=====] - 0s 2ms/step - loss: 0.6249 - accuracy:  
0.6499 - val\_loss: 0.6306 - val\_accuracy: 0.6252  
Epoch 83/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6246 - accuracy:  
0.6500 - val\_loss: 0.6304 - val\_accuracy: 0.6256  
Epoch 84/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6244 - accuracy:  
0.6512 - val\_loss: 0.6306 - val\_accuracy: 0.6249  
Epoch 85/100  
81/81 [=====] - 0s 3ms/step - loss: 0.6243 - accuracy:  
0.6506 - val\_loss: 0.6301 - val\_accuracy: 0.6276  
Epoch 86/100

```

81/81 [=====] - 0s 2ms/step - loss: 0.6241 - accuracy:
0.6520 - val_loss: 0.6309 - val_accuracy: 0.6239
Epoch 87/100
81/81 [=====] - 0s 2ms/step - loss: 0.6242 - accuracy:
0.6499 - val_loss: 0.6304 - val_accuracy: 0.6228
Epoch 88/100
81/81 [=====] - 0s 3ms/step - loss: 0.6239 - accuracy:
0.6510 - val_loss: 0.6298 - val_accuracy: 0.6300
Epoch 89/100
81/81 [=====] - 0s 3ms/step - loss: 0.6239 - accuracy:
0.6527 - val_loss: 0.6305 - val_accuracy: 0.6222
Epoch 90/100
81/81 [=====] - 0s 2ms/step - loss: 0.6237 - accuracy:
0.6508 - val_loss: 0.6303 - val_accuracy: 0.6222
Epoch 91/100
81/81 [=====] - 0s 2ms/step - loss: 0.6235 - accuracy:
0.6512 - val_loss: 0.6300 - val_accuracy: 0.6256
Epoch 92/100
81/81 [=====] - 0s 3ms/step - loss: 0.6234 - accuracy:
0.6521 - val_loss: 0.6296 - val_accuracy: 0.6269
Epoch 93/100
81/81 [=====] - 0s 3ms/step - loss: 0.6233 - accuracy:
0.6514 - val_loss: 0.6295 - val_accuracy: 0.6256
Epoch 94/100
81/81 [=====] - 0s 2ms/step - loss: 0.6232 - accuracy:
0.6504 - val_loss: 0.6296 - val_accuracy: 0.6245
Epoch 95/100
81/81 [=====] - 0s 2ms/step - loss: 0.6230 - accuracy:
0.6524 - val_loss: 0.6295 - val_accuracy: 0.6232
Epoch 96/100
81/81 [=====] - 0s 3ms/step - loss: 0.6229 - accuracy:
0.6523 - val_loss: 0.6294 - val_accuracy: 0.6239
Epoch 97/100
81/81 [=====] - 0s 2ms/step - loss: 0.6228 - accuracy:
0.6518 - val_loss: 0.6294 - val_accuracy: 0.6310
Epoch 98/100
81/81 [=====] - 0s 2ms/step - loss: 0.6227 - accuracy:
0.6520 - val_loss: 0.6294 - val_accuracy: 0.6235
Epoch 99/100
81/81 [=====] - 0s 3ms/step - loss: 0.6225 - accuracy:
0.6531 - val_loss: 0.6303 - val_accuracy: 0.6215
Epoch 100/100
81/81 [=====] - 0s 2ms/step - loss: 0.6225 - accuracy:
0.6525 - val_loss: 0.6292 - val_accuracy: 0.6279
1288/1288 [=====] - 1s 685us/step
92/92 [=====] - 0s 758us/step
82/82 [=====] - 0s 834us/step

```



#### Training Classification Report:

	precision	recall	f1-score	support
0.0	0.653	0.654	0.654	20520
1.0	0.656	0.656	0.656	20682
accuracy			0.655	41202
macro avg	0.655	0.655	0.655	41202
weighted avg	0.655	0.655	0.655	41202

#### Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.639	0.583	0.610	1466
1.0	0.619	0.672	0.645	1477
accuracy			0.628	2943
macro avg	0.629	0.628	0.627	2943
weighted avg	0.629	0.628	0.627	2943

#### Testing Classification Report:

	precision	recall	f1-score	support
0.0	0.641	0.560	0.598	1303
1.0	0.613	0.689	0.649	1314

accuracy			0.625	2617
macro avg	0.627	0.625	0.623	2617
weighted avg	0.627	0.625	0.624	2617

## Decision Trees

```
[ ]: from sklearn.tree import DecisionTreeClassifier

dt_cls = DecisionTreeClassifier(criterion = 'gini', splitter = 'best',
    ↳max_depth = 25, min_samples_split = 2, )
dt_cls.fit(x_train_aug.reshape(n_aug_train, -1), y_train_aug)
preds_train = dt_cls.predict(x_train_aug.reshape(n_aug_train, -1))
preds_val    = dt_cls.predict(x_val.reshape(n_val, -1))
preds_test   = dt_cls.predict(x_test.reshape(n_test, -1))

print("Training Classification Report: \n", classification_report(y_train_aug,
    ↳preds_train, digits = N_DIGITS))
print("\nValidation Classification Report: \n", classification_report(y_val,
    ↳preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test,
    ↳preds_test, digits = N_DIGITS))
```

### Training Classification Report:

	precision	recall	f1-score	support
0.0	0.799	0.998	0.887	20520
1.0	0.997	0.751	0.856	20682
accuracy			0.874	41202
macro avg	0.898	0.874	0.872	41202
weighted avg	0.898	0.874	0.872	41202

### Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.709	0.947	0.811	1466
1.0	0.921	0.614	0.737	1477
accuracy			0.780	2943
macro avg	0.815	0.780	0.774	2943
weighted avg	0.815	0.780	0.774	2943

### Testing Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

	0.0	0.719	0.944	0.816	1303
	1.0	0.920	0.635	0.751	1314
accuracy				0.789	2617
macro avg		0.819	0.789	0.784	2617
weighted avg		0.820	0.789	0.784	2617

## XGBoost

```
[ ]: tf.keras.backend.clear_session()

from xgboost import XGBClassifier

xgb_cls = XGBClassifier(max_depth = 10, objective = 'reg:logistic',
                        num_parallel_tree = 20, booster = 'gbtree',
                        gamma = 0.5, tree_method = 'gpu_hist', subsample = 0.4, reg_lambda = 1)
xgb_cls.fit(x_train_aug.reshape(n_aug_train, -1), y_train_aug)
preds_train = xgb_cls.predict(x_train_aug.reshape(n_aug_train, -1))
preds_val = xgb_cls.predict(x_val.reshape(n_val, -1))
preds_test = xgb_cls.predict(x_test.reshape(n_test, -1))

print("Training Classification Report: \n", classification_report(y_train_aug, preds_train, digits = N_DIGITS))
print("\nValidation Classification Report: \n", classification_report(y_val, preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test, preds_test, digits = N_DIGITS))
```

### Training Classification Report:

	precision	recall	f1-score	support
0.0	0.986	0.999	0.993	20520
1.0	0.999	0.986	0.992	20682
accuracy			0.993	41202
macro avg	0.993	0.993	0.993	41202
weighted avg	0.993	0.993	0.993	41202

### Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.899	0.951	0.924	1466
1.0	0.948	0.894	0.921	1477
accuracy			0.923	2943



macro avg	0.924	0.923	0.922	2943
weighted avg	0.924	0.923	0.922	2943

Testing Classification Report:

	precision	recall	f1-score	support
0.0	0.909	0.954	0.931	1303
1.0	0.952	0.906	0.928	1314
accuracy			0.930	2617
macro avg	0.931	0.930	0.930	2617
weighted avg	0.931	0.930	0.930	2617

## SVM

```
[ ]: from sklearn.svm import SVC

svm_cls = SVC(kernel = 'rbf', max_iter = 250, verbose= True)
svm_cls.fit(x_train_aug.reshape(n_aug_train, -1), y_train_aug)
preds_train = svm_cls.predict(x_train_aug.reshape(n_aug_train, -1))
preds_val   = svm_cls.predict(x_val.reshape(n_val, -1))
preds_test  = svm_cls.predict(x_test.reshape(n_test, -1))

print("Training Classification Report: \n", classification_report(y_train_aug,
↳ preds_train, digits = N_DIGITS))
print("\nValidation Classification Report: \n", classification_report(y_val,
↳ preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test,
↳ preds_test, digits = N_DIGITS))
```

[LibSVM]

```
c:\ProgramData\Anaconda3\envs\ml\lib\site-packages\sklearn\svm\_base.py:301:
ConvergenceWarning: Solver terminated early (max_iter=250). Consider pre-
processing your data with StandardScaler or MinMaxScaler.
warnings.warn(
```

Training Classification Report:

	precision	recall	f1-score	support
0.0	0.437	0.618	0.512	20520
1.0	0.357	0.211	0.265	20682
accuracy			0.413	41202
macro avg	0.397	0.414	0.388	41202
weighted avg	0.397	0.413	0.388	41202

#### Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.432	0.596	0.501	1466
1.0	0.357	0.222	0.274	1477
accuracy			0.408	2943
macro avg	0.394	0.409	0.387	2943
weighted avg	0.394	0.408	0.387	2943

#### Testing Classification Report:

	precision	recall	f1-score	support
0.0	0.423	0.601	0.497	1303
1.0	0.323	0.189	0.238	1314
accuracy			0.394	2617
macro avg	0.373	0.395	0.368	2617
weighted avg	0.373	0.394	0.367	2617

#### CNN Model

```
[ ]: cnn = tf.keras.Sequential([
    tf.keras.layers.Conv2D(64, (3,3), padding = 'same', activation = 'relu',
    ↪input_shape = x_train[0].shape),
    tf.keras.layers.MaxPool2D((3,3), padding = 'same'),

    tf.keras.layers.UpSampling2D((4,4)),

    tf.keras.layers.Conv2D(32, (3,3), padding = 'same', activation = 'relu'),
    tf.keras.layers.MaxPool2D((3,3), padding = 'same'),

    tf.keras.layers.Conv2D(32, (2,2), padding = 'same', activation = 'relu'),
    tf.keras.layers.MaxPool2D((2,2), padding = 'same'),

    tf.keras.layers.Flatten(),

    tf.keras.layers.Dense(512, activation = 'relu'),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(512, activation = 'relu'),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(2, activation = 'sigmoid')
])
```

```

cnn.compile(
    optimizer = tf.keras.optimizers.Adam(learning_rate=0.001),
    loss = 'sparse_categorical_crossentropy',
    metrics = ['accuracy']
)

history = cnn.fit(
    x_train_aug, y_train_aug, batch_size = 512,
    epochs = 100,
    shuffle = True,
    validation_data = [x_val, y_val],
    validation_batch_size = 256,
    callbacks = [
        tf.keras.callbacks.EarlyStopping(monitor = 'val_loss', patience=
↪ 20),
        tf.keras.callbacks.ModelCheckpoint(str = 'val_loss',
↪ save_best_only = True, save_weights_only = True, filepath= "CNN_Aug_HSV")
    ]
)

fig, ax = plot_history(history)

cnn.load_weights('CNN_Aug_HSV')

fig.savefig(PLOT_SAVE + "CNN_aug_losscurve_HSV.png", facecolor = 'white')
preds_train = np.argmax(cnn.predict(x_train_aug), axis = 1)
preds_val    = np.argmax(cnn.predict(x_val), axis = 1)
preds_test   = np.argmax(cnn.predict(x_test), axis = 1)

print("Training Classification Report: \n", classification_report(y_train_aug,
↪ preds_train, digits = N_DIGITS))
print("\nValidation Classification Report: \n", classification_report(y_val,
↪ preds_val, digits = N_DIGITS))
print("\nTesting Classification Report: \n", classification_report(y_test,
↪ preds_test, digits = N_DIGITS))

```

Epoch 1/100

81/81 [=====] - 5s 28ms/step - loss: 0.3102 - accuracy: 0.8585 - val\_loss: 0.0527 - val\_accuracy: 0.9830

Epoch 2/100

81/81 [=====] - 2s 22ms/step - loss: 0.0798 - accuracy: 0.9740 - val\_loss: 0.0355 - val\_accuracy: 0.9901

Epoch 3/100


81/81 [=====] - 2s 21ms/step - loss: 0.0479 - accuracy: 0.9852 - val\_loss: 0.0367 - val\_accuracy: 0.9891

Epoch 4/100

81/81 [=====] - 2s 21ms/step - loss: 0.0421 - accuracy:

0.9870 - val\_loss: 0.0364 - val\_accuracy: 0.9891  
Epoch 5/100  
81/81 [=====] - 2s 22ms/step - loss: 0.0374 - accuracy:  
0.9888 - val\_loss: 0.0265 - val\_accuracy: 0.9925  
Epoch 6/100  
81/81 [=====] - 2s 22ms/step - loss: 0.0363 - accuracy:  
0.9888 - val\_loss: 0.0269 - val\_accuracy: 0.9922  
Epoch 7/100  
81/81 [=====] - 2s 21ms/step - loss: 0.0325 - accuracy:  
0.9907 - val\_loss: 0.0259 - val\_accuracy: 0.9908  
Epoch 8/100  
81/81 [=====] - 2s 20ms/step - loss: 0.0315 - accuracy:  
0.9904 - val\_loss: 0.0272 - val\_accuracy: 0.9925  
Epoch 9/100  
81/81 [=====] - 2s 19ms/step - loss: 0.0287 - accuracy:  
0.9916 - val\_loss: 0.0285 - val\_accuracy: 0.9918  
Epoch 10/100  
81/81 [=====] - 2s 19ms/step - loss: 0.0251 - accuracy:  
0.9920 - val\_loss: 0.0300 - val\_accuracy: 0.9922  
Epoch 11/100  
81/81 [=====] - 2s 21ms/step - loss: 0.0237 - accuracy:  
0.9927 - val\_loss: 0.0295 - val\_accuracy: 0.9915  
Epoch 12/100  
81/81 [=====] - 2s 21ms/step - loss: 0.0224 - accuracy:  
0.9934 - val\_loss: 0.0264 - val\_accuracy: 0.9929  
Epoch 13/100  
81/81 [=====] - 2s 21ms/step - loss: 0.0223 - accuracy:  
0.9931 - val\_loss: 0.0285 - val\_accuracy: 0.9912  
Epoch 14/100  
81/81 [=====] - 2s 21ms/step - loss: 0.0247 - accuracy:  
0.9921 - val\_loss: 0.0276 - val\_accuracy: 0.9915  
Epoch 15/100  
81/81 [=====] - 2s 21ms/step - loss: 0.0244 - accuracy:  
0.9922 - val\_loss: 0.0341 - val\_accuracy: 0.9905  
Epoch 16/100  
81/81 [=====] - 2s 21ms/step - loss: 0.0201 - accuracy:  
0.9938 - val\_loss: 0.0363 - val\_accuracy: 0.9895  
Epoch 17/100  
81/81 [=====] - 2s 21ms/step - loss: 0.0206 - accuracy:  
0.9934 - val\_loss: 0.0292 - val\_accuracy: 0.9925  
Epoch 18/100  
81/81 [=====] - 2s 21ms/step - loss: 0.0172 - accuracy:  
0.9941 - val\_loss: 0.0359 - val\_accuracy: 0.9912  
Epoch 19/100  
81/81 [=====] - 2s 21ms/step - loss: 0.0165 - accuracy:  
0.9950 - val\_loss: 0.0348 - val\_accuracy: 0.9922  
Epoch 20/100  
81/81 [=====] - 2s 21ms/step - loss: 0.0164 - accuracy:

0.9950 - val\_loss: 0.0299 - val\_accuracy: 0.9925  
Epoch 21/100  
81/81 [=====] - 2s 21ms/step - loss: 0.0158 - accuracy:  
0.9953 - val\_loss: 0.0304 - val\_accuracy: 0.9922  
Epoch 22/100  
81/81 [=====] - 2s 21ms/step - loss: 0.0162 - accuracy:  
0.9946 - val\_loss: 0.0322 - val\_accuracy: 0.9915  
Epoch 23/100  
81/81 [=====] - 2s 21ms/step - loss: 0.0157 - accuracy:  
0.9950 - val\_loss: 0.0335 - val\_accuracy: 0.9922  
Epoch 24/100  
81/81 [=====] - 2s 21ms/step - loss: 0.0125 - accuracy:  
0.9958 - val\_loss: 0.0437 - val\_accuracy: 0.9898  
Epoch 25/100  
81/81 [=====] - 2s 21ms/step - loss: 0.0111 - accuracy:  
0.9965 - val\_loss: 0.0377 - val\_accuracy: 0.9925  
Epoch 26/100  
81/81 [=====] - 2s 21ms/step - loss: 0.0120 - accuracy:  
0.9965 - val\_loss: 0.0325 - val\_accuracy: 0.9929  
Epoch 27/100  
81/81 [=====] - 2s 21ms/step - loss: 0.0107 - accuracy:  
0.9967 - val\_loss: 0.0303 - val\_accuracy: 0.9932



1288/1288 [=====] - 2s 1ms/step  
92/92 [=====] - 0s 2ms/step  
82/82 [=====] - 0s 2ms/step  
Training Classification Report:

	precision	recall	f1-score	support
0.0	0.986	0.994	0.990	20520
1.0	0.994	0.986	0.990	20682
accuracy			0.990	41202
macro avg	0.990	0.990	0.990	41202
weighted avg	0.990	0.990	0.990	41202

Validation Classification Report:

	precision	recall	f1-score	support
0.0	0.996	0.986	0.991	1466
1.0	0.986	0.996	0.991	1477
accuracy			0.991	2943
macro avg	0.991	0.991	0.991	2943
weighted avg	0.991	0.991	0.991	2943

Testing Classification Report:

	precision	recall	f1-score	support
0.0	0.996	0.992	0.994	1303
1.0	0.992	0.996	0.994	1314
accuracy			0.994	2617
macro avg	0.994	0.994	0.994	2617
weighted avg	0.994	0.994	0.994	2617